

✓ HW: Sentece contrastive learning

This homework is about learning sentence representation and contrastive learning.

From previous homework, we used to build token/sequence classification task and learn it through only supervised method. In real-world scenario, **human annotation** requires a lot of cost and effort to do. Some annotation tasks might require domain experts such as medical domain, legal domain, etc. However, there are some **unsupervised** methods which are no need any annotations.

Contrastive learning is the popular one of unsupervised learning approach. It will learn the representation via similar and dissimilar examples.

For this homework, we will focus on **SimCSE** framework which is one of contrastive learning techniques. For SimCSE, it will learn sentence embedding by comparing between different views of the same sentence.

In this homework you will perform three main tasks.

1. Train a sentiment classification model using a pretrained model. This model uses freeze weights. That is it treats the pretrained model as a fixed feature extractor.
2. Train a sentiment classification model using a pretrained model. This model also performs weight updates on the base model's weights.
3. Perform SimCSE and use the sentence embedding to perform linear classification.

✓ Install and import libraries

Install the `datasets` library under Huggingface and Pytorch `lightning` framework.

```
1 !pip install -q datasets pytorch-lightning scikit-learn

1 import torch
2 from torch import nn
3 import torch.nn.functional as F
4 from transformers import (
5     AutoTokenizer, AutoModelForSequenceClassification, AutoModel
6 )
7 from datasets import load_dataset
8 import pytorch_lightning as pl
9 from pytorch_lightning import LightningModule, Trainer
10 from torch.utils.data import DataLoader
11 from torchmetrics import Accuracy
12
13 import numpy as np
14 import matplotlib.pyplot as plt
15 from sklearn.manifold import TSNE
```

✓ Setup

The dataset we use for this homework is **Wisesight-Sentiment** ([huggingface](#), [github](#)) dataset. It is a Thai social media dataset which are labeled as **4 classes** e.g. positive, negative, neutral, and question. Furthermore, It contains both Thai, English, Emoji, and etc. That is why we choose the distilled version of multilingual BERT (mBERT) [DistilledBERT paper](#) to be a base model.

```
1 model_name = 'distilbert-base-multilingual-cased'
2 dataset = load_dataset('pythainlp/wisesight_sentiment')
3
4 # Load tokenizer
5 tokenizer = AutoTokenizer.from_pretrained(model_name) # Or a Thai-specific tokenizer if available
```

```
🔗 /usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens),
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(
```

✓ Loading Dataset and DataLoader

✓ Preprocessing step

```
1 # Preprocessing function
2 def preprocess_function(examples):
3     return tokenizer(examples['texts'], padding='max_length', truncation=True)
4
5 # Apply preprocessing
6 encoded_dataset = dataset.map(preprocess_function, batched=True)
7
8 # Change `category` key to `labels`
9 encoded_dataset = encoded_dataset.map(lambda examples: {'labels': [label for label in examples['category']]}, batched=True)
10
```



✓ Define Dataset class

```
1 # Create PyTorch Dataset
2 class SentimentDataset(torch.utils.data.Dataset):
3     def __init__(self, encodings, labels):
4         self.encodings = encodings
5         self.labels = labels
6
7     def __getitem__(self, idx):
8         item = {
9             key: torch.tensor(val) for key, val in self.encodings[idx].items()
10            if key in ['input_ids', 'attention_mask']}
11
12         item['labels'] = torch.tensor(self.labels[idx])
13         return item
14
15     def __len__(self):
16         return len(self.labels)
17
```

✓ Declare Dataset and DataLoader

```
1 # Create Dataset object from DataFrame
2 train_dataset = SentimentDataset(encoded_dataset['train'], encoded_dataset['train']['labels'])
3 val_dataset = SentimentDataset(encoded_dataset['validation'], encoded_dataset['validation']['labels'])
4 test_dataset = SentimentDataset(encoded_dataset['test'], encoded_dataset['test']['labels'])
5
6 # Create dataloaders
7 train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
8 val_loader = DataLoader(val_dataset, batch_size=32)
9 test_loader = DataLoader(test_dataset, batch_size=32)
```

✓ Define base model classes

Here we define model classes which will be used in the next sections.

✓ Base Model class

BaseModel is a parent class for building other models e.g.

- Pretrained LM with a linear classifier
- Fine-tuned LM with a linear classifier
- Contrastive learning based (SimCSE) LM with a linear classifier

```
1 class BaseModel(LightningModule):
2     def __init__(
3         self,
4         model_name: str = 'distilbert-base-multilingual-cased',
5         learning_rate: float = 2e-5
6     ):
7         super().__init__()
8         self.save_hyperparameters()
9
10        self.encoder = AutoModel.from_pretrained(model_name)
11
```

```

11         self.learning_rate = learning_rate
12
13     def get_embeddings(self, input_ids, attention_mask):
14         # TODO 1: get CLS token embedding to represent as a sentence embedding
15         # [CLS] token is on first token
16         return self.encoder(input_ids=input_ids, attention_mask=attention_mask).last_hidden_state[:, 0]
17
18     def configure_optimizers(self):
19         optimizer = torch.optim.AdamW(self.parameters(), lr=self.learning_rate)
20         return optimizer
21
22     def forward(self, input_ids, attention_mask):
23         return self.get_embeddings(input_ids, attention_mask)

```

✓ LMWithLinearClassifier class

LMWithLinearClassifier class is designed to update both LM's parameters in the supervised approach and a linear layer's parameters.

LMWithLinearClassifier consists of

1. ckpt_path (checkpoint path) refers to the best checkpoint after training SimCSE method. We will load the encoder's weights from the checkpoint into the local encoder. This parameter will be in the section of training a linear classifier after SimCSE training part.
2. freeze_weights function is to convert the training status of encoder's weights to non-trainable. This function will be used in the linear classifier training part under both Pretrained LM with a linear classifier and SimCSE with a linear classifier.
3. freeze_encoder_weights is defined to choose whether freeze or unfreeze encoder's weights.

```

1 class LMWithLinearClassifier(BaseModel):
2     def __init__(
3         self,
4         model_name: str = 'distilbert-base-multilingual-cased',
5         ckpt_path: str = None,
6         learning_rate: float = 2e-5,
7         freeze_encoder_weights: bool = False
8     ):
9         super().__init__(
10             model_name,
11             learning_rate
12         )
13         self.save_hyperparameters()
14
15         # TODO 2: load encoder's weights from Pytorch Lightning's checkpoint
16         if ckpt_path is not None:
17             checkpoint = torch.load(ckpt_path)
18             # Load only encoder parts
19             self.encoder.load_state_dict({key.replace("encoder.", ""): val for key, val in checkpoint['state_dict'].items})
20
21         # TODO 3: define a linear classifier which will output the 4 classes
22         self.linear_classifier = nn.Linear(self.encoder.config.hidden_size, 4)
23
24         if freeze_encoder_weights:
25             self.freeze_weights(self.encoder) # Freeze model
26
27         self.accuracy = Accuracy(task='multiclass', num_classes=4)
28
29         # TODO 4: implement `freeze_weights` function which will set requires_grad
30         # in the model.parameters() so that no gradient update will be done on the
31         # base model. Only the linear_layer will be updated.
32         def freeze_weights(self, model):
33             for param in model.parameters():
34                 param.requires_grad = False
35
36         # TODO 5: get logits from the classifier
37         def forward(self, input_ids, attention_mask):
38             embeddings = self.get_embeddings(input_ids, attention_mask)
39             logits = self.linear_classifier(embeddings)
40             return logits
41
42         def training_step(self, batch, batch_idx):
43             # TODO 6.1: implement cross entropy loss for text classification
44             # and log loss and acc
45             logits = self.forward(batch['input_ids'], batch['attention_mask'])
46             loss = F.cross_entropy(logits, batch['labels'])
47
48             acc = self.accuracy(logits, batch['labels'])
49
50             self.log('train_loss', loss, on_step=True, on_epoch=True, prog_bar=True, logger=True)
51             self.log('train_acc', acc, on_step=True, on_epoch=True, prog_bar=True, logger=True)
52
53         return loss

```

```

54
55 def validation_step(self, batch, batch_idx):
56     # TODO 6.2: implement same as `training_step`
57     logits = self.forward(batch['input_ids'], batch['attention_mask'])
58     loss = F.cross_entropy(logits, batch['labels'])
59
60     acc = self.accuracy(logits, batch['labels'])
61
62     self.log('val_loss', loss, on_step=True, on_epoch=True, prog_bar=True, logger=True)
63     self.log('val_acc', acc, on_step=True, on_epoch=True, prog_bar=True, logger=True)
64
65     return loss
66
67 def test_step(self, batch, batch_idx):
68     # TODO 6.3: implement same as `training_step`
69     logits = self.forward(batch['input_ids'], batch['attention_mask'])
70     loss = F.cross_entropy(logits, batch['labels'])
71
72     acc = self.accuracy(logits, batch['labels'])
73
74     self.log('test_loss', loss, on_step=True, on_epoch=True, prog_bar=True, logger=True)
75     self.log('test_acc', acc, on_step=True, on_epoch=True, prog_bar=True, logger=True)
76
77     return loss

```

✓ Pretrained LM with a linear classifier

To benchmark models, we need to have some baselines to compare how good the models' performance are.

The simplest baseline to measure the contrastive learning-based method is the pretrained LM which just fine-tunes only the last linear classifier head to predict sentiments (positive/negative/neutral/questions).

✓ Define model

```

1 pretrained_lm_w_linear_model = LMWithLinearClassifier(
2     model_name,
3     ckpt_path=None,
4     freeze_encoder_weights=True
5 )

```

✓ Train a linear classifier

```

1 # Create a ModelCheckpoint callback (recommended way):
2 pretrained_lm_w_linear_checkpoint_callback = pl.callbacks.ModelCheckpoint(
3     monitor="val_acc", # Metric to monitor
4     mode="max", # "min" for loss, "max" for accuracy
5     save_top_k=1, # Save only the best model(s)
6     save_weights_only=True, # Saves only weights, not the entire model
7     dirpath="./checkpoints/", # Path where the checkpoints will be saved
8     filename="best_pretrained_w_linear_model-{epoch}-{val_acc:.2f}", # Customized name for the checkpoint
9     verbose=True,
10 )
11
12 # Initialize trainer
13 pretrained_lm_w_linear_trainer = Trainer(
14     max_epochs=3,
15     accelerator='auto',
16     callbacks=[pretrained_lm_w_linear_checkpoint_callback], # Add the ModelCheckpoint callback
17     gradient_clip_val=1.0,
18     precision=16, # Mixed precision training
19     devices=1,
20 )
21
22 # Train the model
23 pretrained_lm_w_linear_trainer.fit(pretrained_lm_w_linear_model, train_loader, val_loader)

```

```

INFO:pytorch_lightning.utilities.rank_zero:Using 16bit Automatic Mixed Precision (AMP)
INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used: True
INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:pytorch_lightning.callbacks.model_summary:
  | Name                | Type                | Params | Mode
-----|-----|-----|-----
0 | encoder              | DistilBertModel     | 134 M  | eval
1 | linear_classifier    | Linear              | 3.1 K  | train
2 | accuracy             | MulticlassAccuracy  | 0      | train
-----|-----|-----|-----
3.1 K   Trainable params
134 M   Non-trainable params
134 M   Total params
538.949 Total estimated model params size (MB)
2       Modules in train mode
92      Modules in eval mode

Epoch 2: 100%

676/676 [01:39<00:00, 6.77it/s, v_num=1, train_loss_step=1.150, train_acc_step=0.429, val_loss_step=0.805, val_acc_step=0.500, val_loss_epoch=1.030, val_acc_epoch=0.537]
INFO:pytorch_lightning.utilities.rank_zero:Epoch 0, global step 676: 'val_acc' reached 0.53702 (best 0.53702), saving mo

```

✓ Evaluate

```

1 pretrained_lm_w_linear_result = pretrained_lm_w_linear_trainer.test(pretrained_lm_w_linear_model, test_loader)
2 pretrained_lm_w_linear_result

```

```

INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
Testing DataLoader 0: 100%
84/84 [00:10<00:00, 7.64it/s]

```

Test metric	DataLoader 0
test_acc_epoch	0.5439910292625427
test_loss_epoch	1.0273537635803223

✓ 2) Fine-tuned LM

This is the same as part 1, but you will also gradient update on the base model weights.

✓ Define model

```

1 finetuned_lm_w_linear_model = LMWithLinearClassifier(
2     model_name,
3     ckpt_path=None,
4     freeze_encoder_weights=False
5 )

```

✓ Train both LM and a linear classifier

```

1 # Create a ModelCheckpoint callback (recommended way):
2 finetuned_lm_w_linear_checkpoint_callback = pl.callbacks.ModelCheckpoint(
3     monitor="val_acc", # Metric to monitor
4     mode="max", # "min" for loss, "max" for accuracy
5     save_top_k=1, # Save only the best model(s)
6     save_weights_only=True, # Saves only weights, not the entire model
7     dirpath="./checkpoints/", # Path where the checkpoints will be saved
8     filename="best_finetuned_w_linear_model-{epoch}-{val_acc:.2f}", # Customized name for the checkpoint
9     verbose=True,
10 )
11
12 # Initialize trainer
13 finetuned_lm_w_linear_trainer = Trainer(
14     max_epochs=3,
15     accelerator='auto',
16     callbacks=[finetuned_lm_w_linear_checkpoint_callback], # Add the ModelCheckpoint callback
17     gradient_clip_val=1.0,
18     precision=16, # Mixed precision training
19     devices=1,
20 )

```

```
21
22 # Train the model
23 finetuned_lm_w_linear_trainer.fit(finetuned_lm_w_linear_model, train_loader, val_loader)
```

INFO:pytorch_lightning.utilities.rank_zero:Using 16bit Automatic Mixed Precision (AMP)

INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used: True

INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU cores

INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs

INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

INFO:pytorch_lightning.callbacks.model_summary:

	Name	Type	Params	Mode
0	encoder	DistilBertModel	134 M	eval
1	linear_classifier	Linear	3.1 K	train
2	accuracy	MulticlassAccuracy	0	train

134 M Trainable params

0 Non-trainable params

134 M Total params

538.949 Total estimated model params size (MB)

2 Modules in train mode

92 Modules in eval mode

Epoch 2: 100%

676/676 [04:44<00:00, 2.37it/s, v_num=2, train_loss_step=0.471, train_acc_step=0.857, val_loss_step=0.431, val_acc_step=1.000, val_loss_epoch=0.748, val_acc_epoch=0.67055]

INFO:pytorch_lightning.utilities.rank_zero:Epoch 0, global step 676: 'val_acc' reached 0.67055 (best 0.67055), saving mo

✓ Evaluate

```
1 finetuned_lm_w_linear_result = finetuned_lm_w_linear_trainer.test(finetuned_lm_w_linear_model, test_loader)
2 finetuned_lm_w_linear_result
```

INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

Testing DataLoader 0: 100%

84/84 [00:10<00:00, 7.65it/s]

Test metric	DataLoader 0
test_acc_epoch	0.696368396282196
test_loss_epoch	0.7398968935012817

✓ Contrastive-based model (SimCSE) with a linear classifier

SimCSE (Simple Contrastive Learning of Sentence Embeddings) is a self-supervised learning method that learns high-quality sentence embeddings without relying on any labeled data. It leverages contrastive learning, a technique where similar examples are encouraged to have similar representations, while dissimilar examples are pushed apart in representation space.

Here's the core idea in a nutshell:

- **Data Augmentation:** SimCSE starts with a batch of sentences. For each sentence, it creates two slightly different "views" of the same sentence. These views are created through simple augmentations, like dropout (randomly masking some words) or other minor perturbations. These augmented sentences are semantically similar to the original.
- **Contrastive Objective:** The core of SimCSE is a contrastive loss function. It treats the two different views of the same sentence as a positive pair – the model should learn to make their embeddings similar. All other sentences in the batch (including their augmented versions) are treated as negative pairs – their embeddings should be dissimilar.
- **Learning:** The model is trained to minimize this contrastive loss. This forces the model to learn sentence embeddings that are robust to the augmentations and capture the underlying semantic meaning of the sentences. Sentences with similar meanings will have embeddings close together, while sentences with different meanings will have embeddings far apart.

Paper: <https://arxiv.org/pdf/2104.08821.pdf>

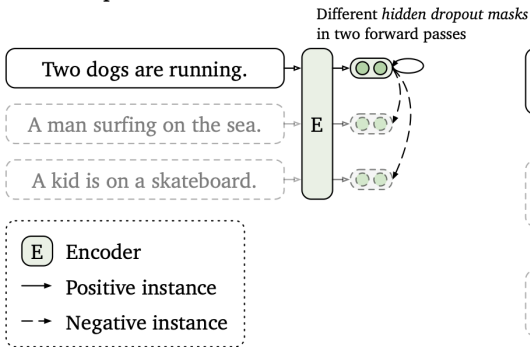
Unsupervised SimCSE is the foundation of the SimCSE method. It's a way to learn sentence embeddings without any labeled data.

Core idea of its concept

- **Dropout as Augmentation:** The key idea in unsupervised SimCSE is to use dropout (randomly masking some words during training) as a form of minimal data augmentation.
- **Two Views:** When you feed the same sentence through your transformer model twice, with dropout turned on, you get two slightly different representations (embeddings) of that sentence. These are like two "views" of the same sentence.

- **Contrastive Learning:** The two embeddings of the same sentence (the "views") are treated as a positive pair. The model is trained to make these embeddings similar to each other. The embeddings of different sentences in the batch are treated as negative pairs. The model is trained to make these embeddings dissimilar to each other.

(a) Unsupervised SimCSE



(b) Supervised SimCSE

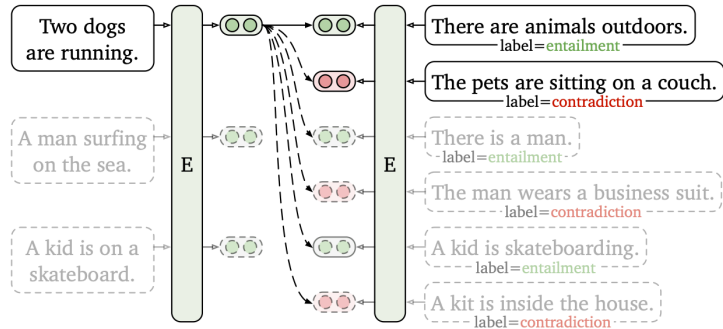


Figure 1: (a) Unsupervised SimCSE predicts the input sentence itself from in-batch negatives, with different hidden dropout masks applied. (b) Supervised SimCSE leverages the NLI datasets and takes the entailment (premise-hypothesis) pairs as positives, and contradiction pairs as well as other in-batch instances as negatives.

✓ Defined Unsupervised SimCSE model and InfoNCE loss

$$L_{UnsupervisedInfoNCE} = -\log \frac{e^{\cos(z_i, z_i)/\tau}}{e^{\cos(z_i, z_i)/\tau} + \sum_{k=0}^N (e^{\cos(z_i, z_k)/\tau})}$$

Notation

z_i indicates the anchor representation (the representation that we are focusing on). The anchor sentence is the initial sentence which its representation is augmented by the dropout masking layer.

z_j indicates the positive representation (the representation that has the same semantic direction). The positive sentence is the same sentence as the anchor one but the positive representation is augmented in different way by the same dropout masking layer.

z_k indicates the negative representation (the representation that has the opposite semantic direction). The negative sentence are the other sentences sampled besides the anchor/positive sentence.

$\cos(\cdot, \cdot)$ is cosine similarity function

N is the number of negative examples

Hint

For loss calculation section, I suggest you to use `F.crossentropy` function and the idea of in-batch negative sampling.

```
Generated code may be subject to a license | clulab/neuralbiocontext
1 class UnsupervisedSimCSE(BaseModel):
2     def __init__(
3         self,
4         model_name: str = 'distilbert-base-multilingual-cased',
5         learning_rate: float = 2e-6,
6         temperature: float = 0.05,
7     ):
8         super().__init__(
9             model_name,
10            learning_rate
11        )
12        self.save_hyperparameters()
13        self.temperature = temperature
14
15        # TODO 7: enable dropout masking in transformer layers to do data augmentation
16        # Dropout layers behave differently during training and inference
17        # https://discuss.pytorch.org/t/if-my-model-has-dropout-do-i-have-to-alternate-between-model-eval-and-model-train-
18        self.dropout = nn.Dropout(p=0.1)
19
20    def forward(self, input_ids, attention_mask):
21        # TODO 8: get sentence embeddings
22        embeddings = self.encoder(input_ids=input_ids, attention_mask=attention_mask).last_hidden_state[:, 0]
23        dropped_out = self.dropout(embeddings)
24        return dropped_out
25
26    def training_step(self, batch, batch_idx):
```

```

27     # TODO 9.1: implement unsupervised InfoNCE loss
28     input_ids = batch['input_ids']
29     attention_mask = batch['attention_mask']
30
31     # First forward pass
32     embeddings1 = self(input_ids, attention_mask)
33
34     # Second forward pass with different dropout
35     embeddings2 = self(input_ids, attention_mask)
36
37     ## Combine embeddings
38     cos_sim = F.cosine_similarity(embeddings1.unsqueeze(1), embeddings2.unsqueeze(0), dim=-1) / self.temperature
39     exp_cos_sim = torch.exp(cos_sim)
40
41     ## Calculate loss
42     exp_of_zi_zj = torch.diag(exp_cos_sim)
43     sum_exp_of_zi_zk = torch.sum(exp_cos_sim, dim=-1) - exp_of_zi_zj
44     loss = -torch.log(exp_of_zi_zj / (exp_of_zi_zj * sum_exp_of_zi_zk)).mean()
45
46     ## Log loss
47     self.log('train_loss', loss, prog_bar=True, on_step=True, on_epoch=True)
48
49     return loss
50
51 def validation_step(self, batch, batch_idx):
52     # TODO 9.2: implement the same as `training_step`
53     input_ids = batch['input_ids']
54     attention_mask = batch['attention_mask']
55
56     # First forward pass
57     embeddings1 = self(input_ids, attention_mask)
58
59     # Second forward pass with different dropout
60     embeddings2 = self(input_ids, attention_mask)
61
62     ## Combine embeddings
63     cos_sim = F.cosine_similarity(embeddings1.unsqueeze(1), embeddings2.unsqueeze(0), dim=-1) / self.temperature
64     exp_cos_sim = torch.exp(cos_sim)
65
66     ## Calculate loss
67     exp_of_zi_zj = torch.diag(exp_cos_sim)
68     sum_exp_of_zi_zk = torch.sum(exp_cos_sim, dim=-1) - exp_of_zi_zj
69     loss = -torch.log(exp_of_zi_zj / (exp_of_zi_zj * sum_exp_of_zi_zk)).mean()
70
71     ## Log loss
72     self.log('val_loss', loss, prog_bar=True, on_step=True, on_epoch=True)
73
74     return loss
75
76 def test_step(self, batch, batch_idx):
77     # TODO 9.3: implement the same as `training_step`
78     input_ids = batch['input_ids']
79     attention_mask = batch['attention_mask']
80
81     # First forward pass
82     embeddings1 = self(input_ids, attention_mask)
83
84     # Second forward pass with different dropout
85     embeddings2 = self(input_ids, attention_mask)
86
87     ## Combine embeddings
88     cos_sim = F.cosine_similarity(embeddings1.unsqueeze(1), embeddings2.unsqueeze(0), dim=-1) / self.temperature
89     exp_cos_sim = torch.exp(cos_sim)
90
91     ## Calculate loss
92     exp_of_zi_zj = torch.diag(exp_cos_sim)
93     sum_exp_of_zi_zk = torch.sum(exp_cos_sim, dim=-1) - exp_of_zi_zj
94     loss = -torch.log(exp_of_zi_zj / (exp_of_zi_zj * sum_exp_of_zi_zk)).mean()
95
96     ## Log loss
97     self.log('test_loss', loss, prog_bar=True, on_step=True, on_epoch=True)
98
99     return loss

```

✓ Train LM through SimCSE approach

```

1 # Initialize model
2 model = UnsupervisedSimCSE()
3
4 # Initialize trainer
5 simcse_trainer = Trainer(

```



```

9 simcse_trainer = Trainer(
10     max_epochs=3,
11     accelerator='auto',
12     devices=1,
13     gradient_clip_val=1.0,
14     precision=16 # Mixed precision training
15 )
16 # Train the model
17 simcse_trainer.fit(model, train_loader)
18
19 # Save the latest checkpoint
20 simcse_trainer.save_checkpoint('/content/latest_simcse_checkpoint.ckpt')

```

```

/usr/local/lib/python3.11/dist-packages/lightning_fabric/connector.py:572: `precision=16` is supported for historical re
INFO:pytorch_lightning.utilities.rank_zero:Using 16bit Automatic Mixed Precision (AMP)
INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used: True
INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs
/usr/local/lib/python3.11/dist-packages/pytorch_lightning/trainer/configuration_validator.py:70: You defined a `validati
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:pytorch_lightning.callbacks.model_summary:
| Name      | Type      | Params | Mode
-----
0 | encoder  | DistilBertModel | 134 M | eval
1 | dropout  | Dropout      | 0     | train
-----
134 M      Trainable params
0          Non-trainable params
134 M      Total params
538.936    Total estimated model params size (MB)
1          Modules in train mode
92         Modules in eval mode

Epoch 2: 100%                                         676/676 [08:30<00:00, 1.32it/s, v_num=8, train_loss_step=4.110, train_loss_epoch=4.310]
INFO:pytorch_lightning.utilities.rank_zero:`Trainer.fit` stopped: `max_epochs=3` reached.

```

Define SimCSE with a linear classifier model

After training SimCSE on the data, we proceed to train a linear classifier on top of the trained model. Be sure to freeze the encoder weights.

```

1 latest_simcse_ckpt_path = '/content/latest_simcse_checkpoint.ckpt'
2
3 simcse_lm_w_linear_model = LMWithLinearClassifier(
4     model_name,
5     ckpt_path=latest_simcse_ckpt_path,
6     freeze_encoder_weights=True
7 )

```

```

<ipython-input-16-396487b92db5>:17: FutureWarning: You are using `torch.load` with `weights_only=False` (the current def
checkpoint = torch.load(ckpt_path)

```

Train a linear classifier

```

1 # Create a ModelCheckpoint callback (recommended way):
2 simcse_lm_w_linear_checkpoint_callback = pl.callbacks.ModelCheckpoint(
3     monitor="val_acc", # Metric to monitor
4     mode="max", # "min" for loss, "max" for accuracy
5     save_top_k=1, # Save only the best model(s)
6     save_weights_only=True, # Saves only weights, not the entire model
7     dirpath="./checkpoints/", # Path where the checkpoints will be saved
8     filename="best_simcse_linear_model-{epoch}-{val_acc:.2f}", # Customized name for the checkpoint
9     verbose=True,
10 )
11
12 # Initialize trainer
13 simcse_lm_w_linear_trainer = Trainer(
14     max_epochs=3,
15     accelerator='auto',
16     callbacks=[simcse_lm_w_linear_checkpoint_callback], # Add the ModelCheckpoint callback
17     gradient_clip_val=1.0,
18     precision=16, # Mixed precision training
19     devices=1,
20 )
21
22 # Train the model
23 simcse_lm_w_linear_trainer.fit(simcse_lm_w_linear_model, train_loader, val_loader)

```

```
➔ /usr/local/lib/python3.11/dist-packages/lightning_fabric/connector.py:572: `precision=16` is supported for historical re
INFO:pytorch_lightning.utilities.rank_zero:Using 16bit Automatic Mixed Precision (AMP)
INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used: True
INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs
/usr/local/lib/python3.11/dist-packages/pytorch_lightning/callbacks/model_checkpoint.py:654: Checkpoint directory /conte
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:pytorch_lightning.callbacks.model_summary:
| Name | Type | Params | Mode
-----
0 | encoder | DistilBertModel | 134 M | eval
1 | linear_classifier | Linear | 3.1 K | train
2 | accuracy | MulticlassAccuracy | 0 | train
-----
3.1 K Trainable params
134 M Non-trainable params
134 M Total params
538.949 Total estimated model params size (MB)
2 Modules in train mode
92 Modules in eval mode
Epoch 2: 100%

676/676 [01:44<00:00, 6.48it/s, v_num=9, train_loss_step=1.270, train_acc_step=0.536, val_loss_step=1.170, val_acc_step=0.500, val_loss_epoch=1.300, val_acc_epoch=0.
INFO:pytorch_lightning.utilities.rank_zero:Epoch 0, global step 676: 'val_acc' reached 0.33694 (best 0.33694), saving mo
```

▼ Evaluate

```
1 simcse_lm_w_linear_result = simcse_lm_w_linear_trainer.test(simcse_lm_w_linear_model, test_loader)
2 simcse_lm_w_linear_result

➔ INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
Testing DataLoader 0: 100%
84/84 [00:11<00:00, 7.59it/s]
```

Test metric	DataLoader 0
test_acc_epoch	0.444777250289917
test_loss_epoch	1.2951160669326782

```
 [{'test_loss_epoch': 1.2951160669326782, 'test_acc_epoch': 0.444777250289917}]
```