

## ✓ Word Tokenizer exercise

In this exercise, you are going to build a set of deep learning models on a (sort of) real world task using pyTorch. PyTorch is a deep learning framework developed by facebook to provide an easier way to use standard layers and networks.

To complete this exercise, you will need to build deep learning models for word tokenization in Thai (ตัดคำภาษาไทย) using NECTEC's BEST corpus. You will build one model for each of the following type:

- Fully Connected (Feedforward) Neural Network
- One-Dimentional Convolution Neural Network (1D-CNN)
- Recurrent Neural Network with Gated Recurrent Unit (GRU)

and one more model of your choice to achieve the highest score possible.

We provide the code for data cleaning and some starter code for PyTorch in this notebook but feel free to modify those parts to suit your needs. Feel free to use additional libraries (e.g. scikit-learn) as long as you have a model for each type mentioned above.

**Don't forget to change hardware accelerator to GPU in Google Colab.**

```
1 %pip install -q wandb torchinfo huggingface_hub lightning
```

↗ Note: you may need to restart the kernel to use updated packages.

```
1 # Run setup code
2 %pip install -q matplotlib pandas tqdm huggingface_hub
3 import os
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 import torch
8 from sklearn.metrics import accuracy_score
9 from huggingface_hub import hf_hub_download
10 from tqdm import tqdm
11
12 %matplotlib inline
13
14 # To guarantee reproducible results
15 torch.manual_seed(5420)
16 torch.backends.cudnn.deterministic = True
17 torch.backends.cudnn.benchmark = False
18 np.random.seed(5420)
```

↗ Note: you may need to restart the kernel to use updated packages.  
/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/tqdm/auto.py:21: TqdmWarning: IPProgress not found. Plea  
from .autonotebook import tqdm as notebook\_tqdm

## ✓ Wandb Setup

We also encourage you to use Wandb which will help you log and visualize your training process.

1. Register [Wandb account](#) (and confirm your email)
2. wandb login and copy paste the API key when prompt

```
1 %env WANDB_API_KEY=648f0ebca50c7021ee306ab62fcbf0029574da
```

↗ env: WANDB\_API\_KEY=648f0ebca50c7021ee306ab62fcbf0029574da

```
1 !wandb login
```

↗ wandb: Currently logged in as: jirayuwat12 (myfistteam). Use `wandb login --relogin` to force relogin

```
1 import wandb
```

```
1 # Check GPU is available
2 torch.cuda.device_count()
```

↗ 0

```
1 # Download dataset
2 hf_hub_download(repo_id="iristun/corpora", filename="corpora.tar.gz", repo_type="dataset", local_dir=".")
```

```
➦ 'corpora.tar.gz'
```

```
1 !tar xvf corpora.tar.gz
```

```
➦ x corpora/  
x corpora/mnist_data/  
x corpora/mnist_data/t10k-images-idx3-ubyte.gz  
x corpora/mnist_data/train-images-idx3-ubyte.gz  
x corpora/mnist_data/.ipynb_checkpoints/  
x corpora/mnist_data/vis_utils.py  
x corpora/mnist_data/__init__.py  
x corpora/mnist_data/load_mnist.py  
x corpora/mnist_data/train-labels-idx1-ubyte.gz  
x corpora/mnist_data/t10k-labels-idx1-ubyte.gz  
x corpora/BEST/  
x corpora/BEST/test/  
x corpora/BEST/test/df_best_article_test.csv  
x corpora/BEST/test/df_best_encyclopedia_test.csv  
x corpora/BEST/test/df_best_novel_test.csv  
x corpora/BEST/test/df_best_news_test.csv  
x corpora/BEST/train/  
x corpora/BEST/train/df_best_encyclopedia_train.csv  
x corpora/BEST/train/df_best_article_train.csv  
x corpora/BEST/train/df_best_news_train.csv  
x corpora/BEST/train/df_best_novel_train.csv  
x corpora/BEST/val/  
x corpora/BEST/val/df_best_encyclopedia_val.csv  
x corpora/BEST/val/df_best_news_val.csv  
x corpora/BEST/val/df_best_article_val.csv  
x corpora/BEST/val/df_best_novel_val.csv  
x corpora/.ipynb_checkpoints/  
x corpora/.ipynb_checkpoints/WordTokenizer.new-checkpoint.ipynb  
x corpora/.ipynb_checkpoints/BackProp-checkpoint.ipynb  
x corpora/.ipynb_checkpoints/WordTokenizer_backup-checkpoint.ipynb  
x corpora/.ipynb_checkpoints/char2vec-checkpoint.ipynb  
x corpora/.ipynb_checkpoints/WordTokenizer-checkpoint.ipynb  
x corpora/cattern/  
x corpora/cattern/gradient_check.py  
x corpora/cattern/.ipynb_checkpoints/  
x corpora/cattern/__init__.py  
x corpora/cattern/data_utils.py  
x corpora/wiki/  
x corpora/wiki/thwiki_chk.txt
```

For simplicity, we are going to build a word tokenization model which is a binary classification model trying to predict whether a character is the beginning of the word or not (if it is, then there is a space in front of it) and without using any knowledge about type of character (vowel, number, English character etc.).

For example,

'แมวตัวน่ารักมาก' -> 'แมว ตัว น่ารัก มาก'

will have these true labels:

[(1,1), (0,0), (1,0) (0,1), ( 0,0), (1,1), (-1,0), (1,0), (1,1), (-1,0), (0,0), (1,1), (1,0), (0,0)]

In this task, we will use only main character you are trying to predict and the characters that surround it (the context) as features. However, you can imagine that a more complex model will try to include more knowledge about each character into the model. You can do that too if you feel like it.

```
1 # Create a character map  
2 CHARS = [  
3     "\n",  
4     " ",  
5     "!",  
6     ":",  
7     "#",  
8     "$",  
9     "%",  
10    "&",  
11    "'",  
12    "(",  
13    ")",  
14    "*",  
15    "+",  
16    ",",  
17    "-",  
18    ".",  
19    "/",  
20    "0",  
21    "1",  
22    "2",  
23    "3",
```

24 "4",  
25 "5",  
26 "6",  
27 "7",  
28 "8",  
29 "9",  
30 ":",  
31 ";",  
32 "<",  
33 "=",  
34 ">",  
35 "?",  
36 "@",  
37 "A",  
38 "B",  
39 "C",  
40 "D",  
41 "E",  
42 "F",  
43 "G",  
44 "H",  
45 "I",  
46 "J",  
47 "K",  
48 "L",  
49 "M",  
50 "N",  
51 "O",  
52 "P",  
53 "Q",  
54 "R",  
55 "S",  
56 "T",  
57 "U",  
58 "V",  
59 "W",  
60 "X",  
61 "Y",  
62 "Z",  
63 "[",  
64 "\\",  
65 "]",  
66 "^",  
67 "\_",  
68 "a",  
69 "b",  
70 "c",  
71 "d",  
72 "e",  
73 "f",  
74 "g",  
75 "h",  
76 "i",  
77 "j",  
78 "k",  
79 "l",  
80 "m",  
81 "n",  
82 "o",  
83 "other",  
84 "p",  
85 "q",  
86 "r",  
87 "s",  
88 "t",  
89 "u",  
90 "v",  
91 "w",  
92 "x",  
93 "y",  
94 "z",  
95 "}",  
96 "~",  
97 "n",  
98 "ñ",  
99 "ü",  
100 "ñ",  
101 "ñ",  
102 "ü",  
103 "ü",  
104 "ä",  
105 "ä",

```

106 "๓",
107 "๔",
108 "๕",
109 "๖",
110 "๗",
111 "๘",
112 "๙",
113 "๐",
114 "๑",
115 "๒",
116 "๓",
117 "๔",
118 "๕",
119 "๖",
120 "๗",
121 "๘",
122 "๙",
123 "๐",
124 "๑",
125 "๒",
126 "๓",
127 "๔",
128 "๕",
129 "๖",
130 "๗",
131 "๘",
132 "๙",
133 "๐",
134 "๑",
135 "๒",
136 "๓",
137 "๔",
138 "๕",
139 "๖",
140 "๗",
141 "๘",
142 "๙",
143 "๐",
144 "๑",
145 "๒",
146 "๓",
147 "๔",
148 "๕",
149 "๖",
150 "๗",
151 "๘",
152 "๙",
153 "๐",
154 "๑",
155 "๒",
156 "๓",
157 "๔",
158 "๕",
159 "๖",
160 "๗",
161 "๘",
162 "๙",
163 "๐",
164 "๑",
165 "๒",
166 "๓",
167 "๔",
168 "๕",
169 "๖",
170 "๗",
171 "๘",
172 "๙",
173 "๐",
174 "๑",
175 "๒",
176 "๓",
177 "๔",
178 "๕",
179 "๖",
180 "\uffeff",
181 ]
182 CHARS_MAP = {v: k for k, v in enumerate(CHARS)}

```

```

1 def create_n_gram_df(df, n_pad):
2     """
3     Given an input dataframe, create a feature dataframe of shifted characters
4     Input:

```

```

5     df: timeseries of size (N)
6     n_pad: the number of context. For a given character at position [idx],
7           character at position [idx-n_pad/2 : idx+n_pad/2] will be used
8           as features for that character.
9
10    Output:
11    dataframe of size (N * n_pad) which each row contains the character,
12          n_pad/2 characters to the left, and n_pad/2 characters to the right
13          of that character.
14    """
15    n_pad_2 = int((n_pad - 1) / 2)
16    for i in range(n_pad_2):
17        df["char-{}".format(i + 1)] = df["char"].shift(i + 1)
18        df["char{}".format(i + 1)] = df["char"].shift(-i - 1)
19    return df[n_pad_2:-n_pad_2]

1 def prepare_feature(best_processed_path, option="train"):
2     """
3     Transform the path to a directory containing processed files
4     into a feature matrix and output array
5     Input:
6     best_processed_path: str, path to a processed version of the BEST dataset
7     option: str, 'train' or 'test'
8     """
9     # we use padding equals 21 here to consider 10 characters to the left
10    # and 10 characters to the right as features for the character in the middle
11    n_pad = 21
12    n_pad_2 = int((n_pad - 1) / 2)
13    pad = [{"char": " ", "target": True}]
14    df_pad = pd.DataFrame(pad * n_pad_2)
15
16    df = []
17    # article types in BEST corpus
18    article_types = ["article", "encyclopedia", "news", "novel"]
19    for article_type in article_types:
20        df.append(
21            pd.read_csv(os.path.join(best_processed_path, option, "df_best_{}_{}.csv".format(article_type, option)))
22        )
23
24    df = pd.concat(df)
25    # pad with empty string feature
26    df = pd.concat((df_pad, df, df_pad))
27
28    # map characters to numbers, use 'other' if not in the predefined character set.
29    df["char"] = df["char"].map(lambda x: CHARS_MAP.get(x, 80))
30
31    # Use nearby characters as features
32    df_with_context = create_n_gram_df(df, n_pad=n_pad)
33
34    char_row = ["char" + str(i + 1) for i in range(n_pad_2)] + ["char-" + str(i + 1) for i in range(n_pad_2)] + ["char"]
35
36    # convert pandas dataframe to numpy array to feed to the model
37    x_char = df_with_context[char_row].to_numpy()
38    y = df_with_context["target"].astype(int).to_numpy()
39
40    return x_char, y

```

Before running the following commands, we must inform you that our data is quite large and loading the whole dataset at once will **use a lot of memory (~6 GB after processing and up to ~12GB while processing)**. We expect you to be running this on Google Cloud or Google Colab so that you will not run into this problem. But, if, for any reason, you have to run this on your PC or machine with not enough memory, you might need to write a data generator to process a few entries at a time then feed it to the model while training.

```

1 # Path to the preprocessed data
2 best_processed_path = "corpora/BEST"

1 # Load preprocessed BEST corpus
2 x_train_char, y_train = prepare_feature(best_processed_path, option="train")
3 x_val_char, y_val = prepare_feature(best_processed_path, option="val")
4 x_test_char, y_test = prepare_feature(best_processed_path, option="test")
5
6 # As a sanity check, we print out the size of the training, val, and test data.
7 print("Training data shape: ", x_train_char.shape)
8 print("Training data labels shape: ", y_train.shape)
9 print("Validation data shape: ", x_val_char.shape)
10 print("Validation data labels shape: ", y_val.shape)
11 print("Test data shape: ", x_test_char.shape)
12 print("Test data labels shape: ", y_test.shape)

```

```

Training data shape: (16461637, 21)
Training data labels shape: (16461637,)
Validation data shape: (2035694, 21)
Validation data labels shape: (2035694,)
Test data shape: (2271932, 21)
Test data labels shape: (2271932,)

```

```

1 # Print some entry from the data to make sure it is the same as what you think.
2 print("First 3 features: ", x_train_char[:3])
3 print("First 30 class labels", y_train[:30])

```

```

First 3 features: [[112. 140. 114. 148. 130. 142. 94. 142. 128. 128. 1. 1. 1. 1.
1. 1. 1. 1. 1. 1. 97.]
[140. 114. 148. 130. 142. 94. 142. 128. 128. 141. 97. 1. 1. 1.
1. 1. 1. 1. 1. 1. 112.]
[114. 148. 130. 142. 94. 142. 128. 128. 141. 109. 112. 97. 1. 1.
1. 1. 1. 1. 1. 1. 140.]]
First 30 class labels [1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0]

```

```

1 # print char of feature 1
2 char = np.array(CHARS)
3
4
5 # A function for displaying our features in text
6 def print_features(tfeature, label, index):
7     feature = np.array(tfeature[index], dtype=int).reshape(21, 1)
8     # Convert to string
9     char_list = char[feature]
10    left = "".join(reversed(char_list[10:20].reshape(10))).replace(" ", "")
11    center = "".join(char_list[20])
12    right = "".join(char_list[0:10].reshape(10)).replace(" ", "")
13    word = "".join([left, " ", center, " ", right])
14    print(center + ": " + word + "\tpred = " + str(label[index]))
15
16
17 for ind in range(0, 30):
18     print_features(x_train_char, y_train, ind)

```

```

ค: ค ณะตุลาการรร pred = 1
ณ: ค ณะตุลาการรร pred = 0
ะ: คณ ะ ตุลาการรร pred = 0
ด: คณ ะ ตุลาการรร pred = 0
ุ: คณ ะ ตุลาการรร pred = 0
ล: คณ ะ ตุลาการรร pred = 0
า: คณ ะ ตุลาการรร pred = 0
ก: คณ ะ ตุลาการรร pred = 0
า: คณ ะ ตุลาการรร pred = 0
ร: คณ ะ ตุลาการรร pred = 0
ร: คณ ะ ตุลาการรร pred = 0
ั: คณ ะ ตุลาการรร pred = 0
ฐ: คณ ะ ตุลาการรร pred = 0
ธ: คณ ะ ตุลาการรร pred = 0
ร: คณ ะ ตุลาการรร pred = 0
ร: คณ ะ ตุลาการรร pred = 0
ม: คณ ะ ตุลาการรร pred = 0
น: คณ ะ ตุลาการรร pred = 0
ู: คณ ะ ตุลาการรร pred = 0
ญ: คณ ะ ตุลาการรร pred = 0
ก: คณ ะ ตุลาการรร pred = 1
ั: คณ ะ ตุลาการรร pred = 0
บ: คณ ะ ตุลาการรร pred = 0
ค: คณ ะ ตุลาการรร pred = 1
ว: คณ ะ ตุลาการรร pred = 0
า: คณ ะ ตุลาการรร pred = 0
ม: คณ ะ ตุลาการรร pred = 0
เ: คณ ะ ตุลาการรร pred = 1
ป: คณ ะ ตุลาการรร pred = 0
ั: คณ ะ ตุลาการรร pred = 0

```

Now, you are going to define the model to be used as your classifier. If you are using Pytorch, please follow the guideline we provide below. You can find more about PyTorch model structure here [documentation](#).

In short, you need to inherit the class `torch.nn.Module` and override the constructor and the method `forward` as shown below:

```

Class Model(torch.nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        #init layer
    def forward(self, x):
        #forward pass the model

```

Also, beware that complex model requires more time to train and your dataset is already quite large. We tested it with a simple 1-hidden-layered feedforward neural network and it used ~5 mins to train 1 epoch.

## ✓ Three-Layer Feedforward Neural Networks

Below, we provide you the code for creating a 3-layer fully connected neural network in PyTorch. This will also serve as the baseline for your other models. Run the code below while making sure you understand what you are doing. Then, report the results.

```
1 import torch.nn.functional as F
2 from torchinfo import summary
3
4
5 class SimpleFeedforwardNN(torch.nn.Module):
6     def __init__(self):
7         super(SimpleFeedforwardNN, self).__init__()
8
9         self.mlp1 = torch.nn.Linear(21, 100)
10        self.mlp2 = torch.nn.Linear(100, 100)
11        self.mlp3 = torch.nn.Linear(100, 100)
12        self.cls_head = torch.nn.Linear(100, 1)
13
14    def forward(self, x):
15        x = F.relu(self.mlp1(x))
16        x = F.relu(self.mlp2(x))
17        x = F.relu(self.mlp3(x))
18        x = self.cls_head(x)
19        out = torch.sigmoid(x)
20        return out
21
22
23 model = SimpleFeedforwardNN() # Initialize model
24 # model.cuda() # specify the location that it is in the GPU
25 model.to('mps:0')
26 summary(model, input_size=(1, 21), device="mps:0") # summarize the model
```

```
→ =====
Layer (type:depth-idx)                   Output Shape          Param #
=====
SimpleFeedforwardNN                     [1, 1]                --
├─Linear: 1-1                           [1, 100]              2,200
├─Linear: 1-2                           [1, 100]              10,100
├─Linear: 1-3                           [1, 100]              10,100
└─Linear: 1-4                           [1, 1]                101
=====
Total params: 22,501
Trainable params: 22,501
Non-trainable params: 0
Total mult-adds (Units.MEGABYTES): 0.02
=====
Input size (MB): 0.00
Forward/backward pass size (MB): 0.00
Params size (MB): 0.09
Estimated Total Size (MB): 0.09
=====
```

## ✓ Test whether the model is working as intended by passing dummy input.

```
1 test_X = torch.tensor(np.zeros((64, 21)), dtype=torch.float).to('mps:0')
2 print(model(test_X).shape)
```

```
→ torch.Size([64, 1])
```

A tensor is very similar to numpy, and many numpy functions has a tensor equivalent.

```
1 example_tensor = torch.arange(6)
2 print(example_tensor.shape)
3
4 # addition and multiplication
5 print(example_tensor * 2 + 1)
6
7 # resize
8 example_tensor = example_tensor.view((2, 3))
9 print(example_tensor)
10
11 example_tensor1 = torch.tensor([[[[1, 2, 3, 4], [5, 6, 7, 8], [9, 10, 11, 12], [13, 14, 15, 16]]]], dtype=torch.float)
12 example_tensor2 = torch.ones_like(example_tensor1)
```

```

13 print(example_tensor1.shape, example_tensor2.shape)
14 print(example_tensor1)
15 print(example_tensor2)
16 print(example_tensor1.matmul(example_tensor2))

```

```

→ torch.Size([6])
tensor([ 1,  3,  5,  7,  9, 11])
tensor([[0, 1, 2],
        [3, 4, 5]])
torch.Size([1, 1, 4, 4]) torch.Size([1, 1, 4, 4])
tensor([[[[ 1.,  2.,  3.,  4.],
           [ 5.,  6.,  7.,  8.],
           [ 9., 10., 11., 12.],
           [13., 14., 15., 16.]]]])
tensor([[[[1., 1., 1., 1.],
           [1., 1., 1., 1.],
           [1., 1., 1., 1.],
           [1., 1., 1., 1.]]]])
tensor([[[[10., 10., 10., 10.],
           [26., 26., 26., 26.],
           [42., 42., 42., 42.],
           [58., 58., 58., 58.]]]])

```

To debug, you can always just try passing variables through individual layers by yourself.

```

1 mlp_test = torch.nn.Linear(21, 3).to('mps:0') # a MLP that has 21 input nodes and 3 output nodes
2 print(x_train_char[:4])
3 print(x_train_char[:4].shape)
4 test_input = torch.tensor(x_train_char[:4], dtype=torch.float).to('mps:0')
5 print(mlp_test(test_input).shape)
6 print(mlp_test(test_input))

```

```

→ [[112. 140. 114. 148. 130. 142.  94. 142. 128. 128.   1.   1.   1.   1.
    1.   1.   1.   1.   1.   1.  97.]
 [140. 114. 148. 130. 142.  94. 142. 128. 128. 141.  97.   1.   1.   1.
    1.   1.   1.   1.   1.   1. 112.]
 [114. 148. 130. 142.  94. 142. 128. 128. 141. 109. 112.  97.   1.   1.
    1.   1.   1.   1.   1.   1. 140.]
 [148. 130. 142.  94. 142. 128. 128. 141. 109. 117. 140. 112.  97.   1.
    1.   1.   1.   1.   1.   1. 114.]]
(4, 21)
torch.Size([4, 3])
tensor([[-10.7609,  35.0017, -75.6997],
        [-27.3163,  26.4542, -69.2510],
        [-9.0613,  42.3437, -91.1748],
        [-21.3748,  35.3703, -57.5476]], device='mps:0',
        grad_fn=<LinearBackward0>)

```

## Typical PyTorch training loop

Before the training loop begins, a data loader responsible for generating data in a trainable format has to be created first. In PyTorch, `torch.utils.data.DataLoader` is a readily available class that are commonly used for data preparation. The dataloader takes the object `torch.utils.data.Dataset` as an input. An example of a data loader for this task is shown below. You can read more about the class `Dataset` here [https://pytorch.org/tutorials/beginner/basics/data\\_tutorial.html](https://pytorch.org/tutorials/beginner/basics/data_tutorial.html).

### ✓ Converting the data into trainable format

In order to train the model using the PyTorch frame, the data has to be converted into `Tensor` type. In the cell below, we convert the data into `cuda.FloatTensor` type. You can read more about `Tensor` data type here : <https://pytorch.org/docs/stable/tensors.html>.

```

1 class Dataset(torch.utils.data.Dataset):
2     "Characterizes a dataset for PyTorch"
3
4     def __init__(self, X, Y, dtype="float"):
5         "Initialization"
6         self.X = X
7         self.Y = Y.reshape(-1, 1)
8         if dtype == "float":
9             self.X = torch.tensor(self.X, dtype=torch.float).to('mps:0')
10        elif dtype == "long":
11            self.X = torch.tensor(self.X, dtype=torch.long).to('mps:0')
12            self.Y = torch.tensor(self.Y, dtype=torch.float).to('mps:0')
13
14        def __len__(self):
15            "Denotes the total number of samples"
16            return len(self.X)
17
18        def __getitem__(self, index):

```



```

19         "Generates one sample of data"
20         # Select sample
21         x = self.X[index]
22         y = self.Y[index, :]
23         return x, y

```

In the block below, we initialized the hyperparameters used for the training process. Normally, the optimizer, objective function, and training schedule is initialized here.

```

1 from torch.utils.data import DataLoader
2 import torch.optim as optim
3
4 # hyperparameter initialization
5 NUM_EPOCHS = 3
6 criterion = torch.nn.BCELoss(reduction="none")
7 BATCHS_SIZE = 512
8 optimizer_class = optim.Adam
9 optimizer_params = {"lr": 5e-4}
10
11 config = {
12     "architecture": "simpleff",
13     "epochs": NUM_EPOCHS,
14     "batch_size": BATCHS_SIZE,
15     "optimizer_params": optimizer_params,
16 }
17
18 train_loader = DataLoader(Dataset(x_train_char, y_train, dtype="float"), batch_size=BATCHS_SIZE)
19 val_loader = DataLoader(Dataset(x_val_char, y_val, dtype="float"), batch_size=BATCHS_SIZE)
20 test_loader = DataLoader(Dataset(x_test_char, y_test, dtype="float"), batch_size=BATCHS_SIZE)

```

## ✓ Pytorch Lightning Module

In most of our labs, we will use Pytorch Lightning. PyTorch Lightning is an open-source Python library that provides a high-level interface for PyTorch, making it easier/faster to use. It is considered an industry standard and is used widely on recent huggingface tutorials. Pytorch Lightning makes scaling training of deep learning models simple and hardware agnostic.

If you are not familiar with Lightning, you are encouraged to study from this simple [tutorial](#).

```

1 import pytorch_lightning as pl
2 from pytorch_lightning.callbacks import ModelCheckpoint
3 from torchmetrics.functional import accuracy
4
5
6 class LightningModel(pl.LightningModule):
7     def __init__(
8         self,
9         model=SimpleFeedforwardNN(),
10         criterion=criterion,
11         optimizer_class=optim.Adam,
12         optimizer_params={"lr": 5e-4},
13     ):
14         super().__init__()
15         self.model = model
16         self.criterion = criterion
17         self.optimizer_class = optimizer_class
18         self.optimizer_params = optimizer_params
19
20     def forward(self, x):
21         return self.model(x)
22
23     def training_step(self, batch, batch_idx):
24         X_train, Y_train = batch
25         Y_pred = self.model(X_train)
26         loss = self.criterion(Y_pred, Y_train).mean()
27         self.log("train_loss", loss, on_step=True, on_epoch=True)
28         return loss
29
30     def validation_step(self, batch, batch_idx):
31         X_val, Y_val = batch
32         Y_pred = self.model(X_val)
33         loss = self.criterion(Y_pred, Y_val).mean()
34         self.log("val_loss", loss, on_step=False, on_epoch=True)
35
36         # Convert probabilities to classes.
37         val_pred = (Y_pred >= 0.5).float()
38
39         # Calculate accuracy.

```

```

40     val_acc = accuracy(val_pred, Y_val, task="binary")
41
42     self.log("val_accuracy", val_acc, on_step=False, on_epoch=True)
43     return {"val_loss": loss, "val_accuracy": val_acc}
44
45 def configure_optimizers(self):
46     return self.optimizer_class(self.parameters(), **self.optimizer_params)

```

## Initialize LightningModel and Trainer

```

1 # Initialize LightningModel.
2 lightning_model = LightningModel(
3     model,
4     criterion,
5     optimizer_class,
6     optimizer_params,
7 )
8 # Define checkpoint.
9 feedforward_nn_checkpoint = ModelCheckpoint(
10     monitor="val_accuracy", mode="max", save_top_k=1, dirpath="./checkpoints", filename="feedforward_nn"
11 )
12 # Initialize Trainer
13 trainer = pl.Trainer(
14     max_epochs=NUM_EPOCHS,
15     logger=pl.loggers.WandbLogger(),
16     callbacks=[feedforward_nn_checkpoint],
17     accelerator="mps",
18     devices=1,
19 )

```

GPU available: True (mps), used: True  
 TPU available: False, using: 0 TPU cores  
 HPU available: False, using: 0 HPUs

## Starting the training loop

```

1 # Initialize wandb to log the losses from each step.
2 wandb.init(
3     project="simpleff",
4     config=config,
5 )
6 # Fit model.
7 # trainer.fit(lightning_model, train_loader, val_loader)
8
9 print(f"Best model is saved at {feedforward_nn_checkpoint.best_model_path}")

```

wandb: Currently logged in as: jirayuwat12 (myfistteam). Use `wandb login --relogin` to force relogin  
 wandb: Using wandb-core as the SDK backend. Please refer to <https://wandb.me/wandb-core> for more information.  
 Tracking run with wandb version 0.19.2  
 Run data is saved locally in /Users/jirayuwat/Desktop/2110572-NLP/homework1/wandb/run-20250112\_212353-6nnh4pt9  
 Syncing run [iconic-snow-4](https://wandb.ai/myfistteam/simpleff) to [Weights & Biases \(docs\)](https://wandb.ai/myfistteam/simpleff)  
 View project at <https://wandb.ai/myfistteam/simpleff>  
 View run at <https://wandb.ai/myfistteam/simpleff/runs/6nnh4pt9>  
 Best model is saved at

## Evaluate the model performance on the test set

```

1 from sklearn.metrics import f1_score, precision_score, recall_score
2
3
4 #####
5 # A function to evaluate your model. This function must take test dataloader #
6 # and the input model to return f-score, precision, and recall of the model. #
7 #####
8 def evaluate(test_loader, model):
9     """
10     Evaluate model on the splitted 10 percent testing set.
11     """
12     model.eval()
13     with torch.no_grad():
14         test_loss = []
15         test_pred = []
16         test_true = []
17         for X_test, Y_test in tqdm(test_loader):
18             Y_pred = model(X_test)
19             loss = criterion(Y_pred, Y_test)

```

```

20         test_loss.append(loss)
21         test_pred.append(Y_pred)
22         test_true.append(Y_test)
23
24     avg_test_loss = torch.cat(test_loss, axis=0).mean().item()
25     test_pred = torch.cat(test_pred, axis=0).cpu().detach().numpy()
26     test_true = torch.cat(test_true, axis=0).cpu().detach().numpy()
27
28     prob_to_class = lambda p: 1 if p[0] >= 0.5 else 0
29     test_pred = np.apply_along_axis(prob_to_class, 1, test_pred)
30
31     acc = accuracy_score(test_true, test_pred)
32     f1score = f1_score(test_true, test_pred)
33     precision = precision_score(test_true, test_pred)
34     recall = recall_score(test_true, test_pred)
35
36     return {"accuracy": acc, "f1_score": f1score, "precision": precision, "recall": recall}

1 # Load best model and evaluate it.
2 best_model_path = './checkpoints/feedforward_nn.ckpt'
3 # best_model_path = ... # Insert if you have already trained this model.
4 best_model = LightningModel.load_from_checkpoint(best_model_path, model=SimpleFeedforwardNN())
5 result = evaluate(test_loader, best_model)
6
7 wandb.finish()
8 print(result)

```

 100%|██████████| 4438/4438 [00:34<00:00, 130.50it/s]  
 View run **iconic-snow-4** at: <https://wandb.ai/myfistteam/simpleff/runs/6nnh4pt9>  
 View project at: <https://wandb.ai/myfistteam/simpleff>  
 Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)  
 Find logs at: ./wandb/run-20250112\_212353-6nnh4pt9/logs

## ▼ Debugging


In order to understand what is going on in your model and where the error is, you should try looking at the inputs your model made wrong predictions.

In this task, write a function to print the characters on test data that got wrong prediction along with its context of size 10 (from [x-10] to [x+10]). Examine a few of those and write your assumption on where the model got wrong prediction.

```

1 # TODO#1
2 # Write code to show a few of the errors the models made.
3
4 def show_errors(test_loader: DataLoader, model: LightningModel):
5     """Show errors that the model made"""
6     model.eval()
7
8     # Get the predictions
9     with torch.no_grad():
10         for X_test, Y_test in tqdm(test_loader):
11             # Get the predictions
12             y_pred = model(X_test)
13             y_pred = (y_pred >= 0.5).float()
14             # Convert indices to characters
15             X_text = X_test.cpu().detach().numpy()
16             X_text = np.vectorize(lambda x: CHARS[int(x)])(X_text)
17             X_text = np.apply_along_axis(lambda x: "".join(x), 1, X_text)
18             # Iterate over the batch
19             for index, (text, pred, true) in enumerate(zip(X_text, y_pred, Y_test)):
20                 # Check if the prediction is correct
21                 if pred != true:
22                     show_text = '|' + text[0] + '|' + text[1:10]
23                     show_text = ''.join([X_text[index-i][0] for i in range(10, 0, -1) if index-i >= 0]) + show_text
24                     error_type = 'FP' if pred == 1 else 'FN'
25                     print(f'{error_type} : {show_text}\n')
26             break
27
28 show_errors(test_loader, best_model)

```

 0%| | 0/4438 [00:00<?, ?it/s] FP : มิ|ร|ปการศึกษ  
 FP : มิ|ร|ปการศึกษ : ม  
 FN : น|ท|ศ|น|แ|ละ|บ|ร|บ|ท|ล|ง|ค|ม|ไ  
 FN : แ|ละ|บ|ร|บ|ท|ล|ง|ค|ม|ไ|ท|ย|The  
 FP : บ|ร|บ|ท|ล|ง|ค|ม|ไ|ท|ย|The Re  
 FN : ร|บ|ท|ล|ง|ค|ม|ไ|ท|ย|The Refo  
 FP : บ|ร|บ|ท|ล|ง|ค|ม|ไ|ท|ย|The Refo  
 FP : P|e|r|s|p|e|c|t|i|v|e|e|r|e|v|e|w|e|r|

FN : rspectiven|ร|ะบวนทัศน์  
 FN : ivenระบวนท|ั|ทัศน์และวิธี  
 FP : และวิธีคิ|ด|แบบแยกส่ว  
 FP : วิธีคิดแบบ|แยกส่ว ลด  
 FN : วิธีคิดแบบ|แยกส่ว ลดส  
 FN : คิดแบบแยกส|่ว ลดส่ว  
 FN : ยกส่ว ลดส|่ว น ได้ทำใ  
 FN : "การศึกษา|ร|ียนรู้"ใน  
 FP : "ษาเรียนรู้|ั|ใน หลาย  
 FP : ษาเรียนรู้|ั|ใน หลายทศ  
 FN : "ใน หลายทศ|วรรษที่ผ่  
 FN : รัชที่ผ่านม|า|กลายเป็  
 FN : องของนักท|า|ชาการด้าน  
 FP : ารด้านศึกษา|าศาสตร์ คร  
 FP : ครูอาจารย์|ั|กระทรวงศึ  
 FP : ารยั กระทรวง|ว|งศึกษาอีก  
 FP : กระทรวงศึษา|า|ธิการ ทบ  
 FP : กระทรวงศึษา|ธ|การ ทบวง  
 FP : วงศึษาอีก|า|ร ทบวงมหา  
 FP : ิการ ทบวงม|ห|วิทยาลัย  
 FP : วย่างต่อเ|น|ื่องยาวนาน  
 FP : เนื่องยาวนาน|าน (เหมือน  
 FN : ที่เรื่องส|ั|ภาพเป็นเ  
 FP : เรื่องสขภ|า|พเป็นเร  
 FN : งแพทย์และ|ร|ังพยาบาล)  
 FN : หยัและโรงพย|า|บาล) การ  
 FN : ดการศึกษา|า|ยได้กระบว  
 FP : ายได้กระบว|น|ทัศน์และว  
 FP : และวิธีคิ|ด|แบบดังกล  
 FN : ละวิธีคิดแ|บ|บดังกล่ว  
 FN : บดังกล่วข|อ|งรัฐ ได้ถ  
 FN : งกล่วของร|ั|รัฐ ได้ถูก  
 FP : รัฐ ได้ถูก|า|พาक्षวิจ  
 FN : รัฐ ได้ถูก|า|พาक्षวิจ  
 FP : ูกวิพาक्षวิ|า|จารย์และด  
 FP : วิพาक्षวิจ|า|รณ์และดกเ  
 FN : ิจารย์และด|ก|เป็นจำเลย  
 FN : ารณ์และดก|ป|เป็นจำเลยจา  
 FN : เป็นจำเลยจ|า|กวิฤตการ  
 FN : นจำเลยจากว|า|กฤตการณท  
 FP : ุณททางสังค|ม|มากมาย อั  
 FN : ทางสังคมม|า|กมาย อันส

## Write your answer here

**Your answer:** Mixed Thai-English text increases complexity, as the model struggles to discern where tokens begin or end. Insufficient training data for Thai-specific patterns and rare sequences further exacerbates the issue. The model may also misinterpret visually similar characters or fail to utilize context effectively, leading to incorrect splits. These factors combined highlight the need for better training datasets and improved handling of Thai linguistic features.

## ✓ Dropout

You might notice that the 3-layered feedforward does not use dropout at all. Now, try adding dropout to the model, run, and report the result again.

```

1 #####
2 # TODO#3: #
3 # Write a model class that return feedforward model with dropout. #
4 #####
5
6
7 class SimpleFeedforwardNNWDropout(torch.nn.Module):
8     def __init__(self):
9         super(SimpleFeedforwardNNWDropout, self).__init__()
10
11         self.mlp1 = torch.nn.Linear(21, 100)
12         self.mlp2 = torch.nn.Linear(100, 100)
13         self.mlp3 = torch.nn.Linear(100, 100)
14         self.cls_head = torch.nn.Linear(100, 1)
15         self.dropout = torch.nn.Dropout(0.5)
16
17     def forward(self, x):
18         x = F.relu(self.mlp1(x))
19         x = self.dropout(x)
20         x = F.relu(self.mlp2(x))
21         x = self.dropout(x)
22         x = F.relu(self.mlp3(x))
23         x = self.dropout(x)


```

```

24         x = self.cls_head(x)
25         out = torch.sigmoid(x)
26         return out

1 #####
2 # TODO#4: #
3 # Write code that performs a training process. Select your batch size carefully#
4 # as it will affect your model's ability to converge and #
5 # time needed for one epoch. #
6 #####
7 # Complete the code to train your model with dropout
8 model_nn_with_dropout = SimpleFeedforwardNNWDropout().to('mps:0')
9 summary(model_nn_with_dropout, input_size=(64, 21), device="mps:0") # summarize the model
10
11 #####
12 # WRITE YOUR CODE BELOW #
13 #####
14 # hyperparameter initialization
15 lightning_model_dropout = LightningModel(
16     model_nn_with_dropout,
17     criterion,
18     optimizer_class,
19     optimizer_params,
20 )
21 # Define checkpoint.
22 feedforward_nn_dropout_checkpoint = ModelCheckpoint(
23     monitor="val_accuracy", mode="max", save_top_k=1, dirpath="./checkpoints", filename="feedforward_nn_dropout"
24 )
25 # Initialize Trainer
26 trainer_dropout = pl.Trainer(
27     max_epochs=NUM_EPOCHS,
28     logger=pl.loggers.WandbLogger(),
29     callbacks=[feedforward_nn_dropout_checkpoint],
30     accelerator="mps",
31     devices=1,
32 )
33 # Initialize wandb to log the losses from each step.
34 wandb.init(
35     project="simpleff_dropout",
36     config=config,
37 )
38 # Fit model.
39 # trainer_dropout.fit(lightning_model_dropout, train_loader, val_loader)
40 lightning_model_dropout = LightningModel.load_from_checkpoint(
41     'feedforward_nn_dropout.ckpt', model=SimpleFeedforwardNNWDropout()
42 )

```

 GPU available: True (mps), used: True  
 TPU available: False, using: 0 TPU cores  
 HPU available: False, using: 0 HPUs  
 Tracking run with wandb version 0.19.2  
 Run data is saved locally in /Users/jirayuwat/Desktop/2110572-NLP/homework1/wandb/run-20250112\_212433-5ftb04og  
 Syncing run [sweet-forest-3](https://wandb.ai/myfistteam/sweet-forest-3) to [Weights & Biases \(docs\)](https://wandb.ai/myfistteam/sweet-forest-3)  
 View project at <https://wandb.ai/myfistteam/sweet-forest-3>  
 View run at <https://wandb.ai/myfistteam/sweet-forest-3/runs/5ftb04og>  
 /Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/pytorch\_lightning/loggers/wandb.py:397: There is a wand
 /Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/pytorch\_lightning/callbacks/model\_checkpoint.py:654: Ch

	Name	Type	Params	Mode
0	model	SimpleFeedforwardNNWDropout	22.5 K	train
1	criterion	BCELoss	0	eval

22.5 K	Trainable params
0	Non-trainable params
22.5 K	Total params
0.090	Total estimated model params size (MB)
6	Modules in train mode
1	Modules in eval mode

/Users/jirayuwat/anaconda3/envs/nlp/lib/pytho  
 /Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/pytorch\_lightning/trainer/connectors/data\_connector.py:  
 Epoch 2: 100%|██████████| 32152/32152 [08:22<00:00, 63.92it/s, v\_num=04og]`Trainer.fit` stopped: `max\_epochs=3` reached.  
 Epoch 2: 100%|██████████| 32152/32152 [08:23<00:00, 63.92it/s, v\_num=04og]

```

1 result = evaluate(test_loader, model_nn_with_dropout.to('mps:0'))
2 print(result)

```

 100%|██████████| 4438/4438 [00:36<00:00, 121.56it/s]  
 {'accuracy': 0.8220316453133281, 'f1\_score': 0.5874255884051582, 'precision': 0.8295794822135185, 'recall': 0.4546991130


## ✓ Convolution Neural Networks

Now, you are going to implement your own 1d-convolution neural networks with the following structure: input -> embedding layer (size 32) -> 1D-convolution layer (100 filters of size 5, strides of 1) -> Dense size 5 (applied across time dimension) -> fully-connected layer (size 100) -> output.

These parameters are simple guidelines to save your time. You can play with them in the final section.

The results should be better than the feedforward model.

Embedding layers turn the input from a one-hot vector into better representations via some feature transform (a simple matrix multiply in this case).



Note you need to flatten the tensor before the final fully connected layer because of dimension mis-match. The tensor could be reshaped using the view method.

Do consult PyTorch documentation on how to use [embedding layers](#) and [1D-cnn](#).

Hint: to apply dense5 across the time dimension you should read about how the multiplication in the dense layer is applied. The output of the 1D-cnn should be [batch x nfilter x sequence length]. We want to apply the dense5 (a weight matrix of size 100 x 5) by multiplying the same set of numbers over the nfilter dimension repeated over the sequence length (this can be possible via broadcasting) which should give an output of [batch x 5 x sequence length]. You might want to use the function [transpose](#) somehow.

Even more hints: <https://stackoverflow.com/questions/58587057/multi-dimensional-inputs-in-pytorch-linear-method>

```
1 #####
2 # TODO#5: #
3 # Write a function that returns convolution nueral network model. #
4 # You can choose any normalization methods, activation function, as well as #
5 # any hyperparameter the way you want. Your goal is to predict a score #
6 # between [0,1] for each input whether it is the beginning of the word or not. #
7 # #
8 # Hint: You should read PyTorch documentation to see the list of available #
9 # layers and options you can use. #
10 #####
11 from torch import nn
12
13 class SimpleCNN(torch.nn.Module):
```

```

14 def __init__(self):
15     super(SimpleCNN, self).__init__()
16     self.embed = torch.nn.Embedding(num_embeddings=178, embedding_dim=32)
17     self.cnn = torch.nn.Conv1d(in_channels=32, out_channels=100, kernel_size=5, padding=2)
18     self.dense5 = torch.nn.Linear(100, 5)
19     self.mlp = torch.nn.Linear(105, 100)
20     self.cls_head = torch.nn.Linear(100, 1)
21
22 def forward(self, x):
23     x = x.long()
24     x = self.embed(x)
25     x = x.transpose(1, 2)
26     x = F.relu(self.cnn(x))
27
28     x = x.transpose(1, 2)
29     x = F.relu(self.dense5(x))
30     x = x.flatten(1)
31
32     x = F.relu(self.mlp(x))
33     x = self.cls_head(x)
34     out = torch.sigmoid(x)
35
36     return out

```

```

1 #####
2 # TODO#6: #
3 # Write code that performs a training process. Select your batch size carefully#
4 # as it will affect your model's ability to converge and #
5 # time needed for one epoch. #
6 #####
7 model_conv1d_nn = SimpleCNN().to('mps:0')
8 summary(model_conv1d_nn, input_size=(64, 21), device="mps:0") # summarize the model
9 #####
10 # WRITE YOUR CODE BELOW #
11 #####
12 # hyperparameter initialization
13 lightning_model_conv1d = LightningModel(
14     model_conv1d_nn,
15     criterion,
16     optimizer_class,
17     optimizer_params,
18 )
19 # Define checkpoint.
20 conv1d_nn_checkpoint = ModelCheckpoint(
21     monitor="val_accuracy", mode="max", save_top_k=1, dirpath="./checkpoints", filename="conv1d_nn"
22 )
23 # Initialize Trainer
24 trainer_conv1d = pl.Trainer(
25     max_epochs=NUM_EPOCHS,
26     logger=pl.loggers.WandbLogger(),
27     callbacks=[conv1d_nn_checkpoint],
28     accelerator="mps",
29     devices=1,
30 )
31 # Initialize wandb to log the losses from each step.
32 wandb.init(
33     project="conv1d_nn",
34     config=config,
35 )
36 # Fit model.
37 trainer_conv1d.fit(lightning_model_conv1d, train_loader, val_loader)

```

```

GPU available: True (mps), used: True
TPU available: False, using: 0 TPU cores
HPU available: False, using: 0 HPUs
/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/pytorch_lightning/loggers/wandb.py:397: There is a wand
/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/pytorch_lightning/callbacks/model_checkpoint.py:654: Ch

```

	Name	Type	Params	Mode
0	model	SimpleCNN	33.0 K	train
1	criterion	BCELoss	0	eval

---

```

33.0 K Trainable params
0 Non-trainable params
33.0 K Total params
0.132 Total estimated model params size (MB)
6 Modules in train mode
1 Modules in eval mode

```

```

/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/pytorch_lightning/trainer/connectors/data_connector.py:
Epoch 2: 100%|██████████| 32152/32152 [09:01<00:00, 59.33it/s, v_num=04og]\Trainer.fit` stopped: `max_epochs=3` reached.
Epoch 2: 100%|██████████| 32152/32152 [09:01<00:00, 59.33it/s, v_num=04og]

```

```
1 result = evaluate(test_loader, model_conv1d_nn.to('mps:0'))
2 print(result)
```

```
100%|██████████| 4438/4438 [00:38<00:00, 114.81it/s]
{'accuracy': 0.9718314632656259, 'f1_score': 0.9502210212480642, 'precision': 0.9359999877409291, 'recall': 0.9648808536}
```

## ✓ Final Section

### PyTorch playground

Now, train the best model you can do for this task. You can use any model structure and function available. Remember that training time increases with the complexity of the model. You might find wandb helpful in tuning of complicated models.

Your model should be better than your CNN or GRU model in the previous sections.

Some ideas to try

1. Tune the parameters
2. Recurrent models
3. CNN-GRU model
4. Improve the learning rate scheduling

```
1 #####
2 # TODO#7 #
3 # Write a class that returns your best model. You can use anything #
4 # you want. The goal here is to create the best model you can think of. #
5 # Your model should get f-score more than 97% from calling evaluate(). #
6 # #
7 # Hint: You should read PyTorch documentation to see the list of available #
8 # layers and options you can use. #
9 #####
10
11
12 class BestModel(torch.nn.Module):
13     def __init__(self):
14         super(BestModel, self).__init__()
15         self.embed = torch.nn.Embedding(num_embeddings=178, embedding_dim=64)
16         self.cnn = torch.nn.Conv1d(in_channels=64, out_channels=100, kernel_size=5, padding=2)
17         self.dense5 = torch.nn.Linear(100, 8)
18         self.mlp = torch.nn.Linear(168, 100)
19         self.cls_head = torch.nn.Linear(100, 1)
20
21     def forward(self, x):
22         x = x.long()
23         x = self.embed(x)
24         x = x.transpose(1, 2)
25         x = F.relu(self.cnn(x))
26
27         x = x.transpose(1, 2)
28         x = F.relu(self.dense5(x))
29         x = x.flatten(1)
30
31         x = F.relu(self.mlp(x))
32         x = self.cls_head(x)
33         out = torch.sigmoid(x)
34
35     return out
36
37
38 #####
39 # TODO#8 #
40 # Write code that perform a trainin loop on this dataset. Select your #
41 # batch size carefully as it will affect your model's ability to converge and #
42 # time needed for one epoch. #
43 # #
44 # #####
45 print("start training")
46 my_best_model = BestModel().to('mps:0')
47 #####
48 # #
49 # WRITE YOUR CODE BELOW #
50 #####
51 # hyperparameter initialization
```



```

14 lightning_model_best = LightningModel(
15     my_best_model,
16     criterion,
17     optimizer_class,
18     optimizer_params,
19 )
20 # Define checkpoint.
21 best_checkpoint = ModelCheckpoint(
22     monitor="val_accuracy", mode="max", save_top_k=1, dirpath="./checkpoints", filename="best_model"
23 )
24 # Initialize Trainer
25 trainer_best = pl.Trainer(
26     max_epochs=NUM_EPOCHS,
27     logger=pl.loggers.WandbLogger(),
28     callbacks=[best_checkpoint],
29     accelerator="mps",
30     devices=1,
31 )
32 # Initialize wandb to log the losses from each step.
33 wandb.init(
34     project="best_model",
35     config=config,
36 )
37 # Fit model.
38 trainer_best.fit(lightning_model_best, train_loader, val_loader)
39

```

```

GPU available: True (mps), used: True
TPU available: False, using: 0 TPU cores
HPU available: False, using: 0 HPUs
/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/pytorch_lightning/loggers/wandb.py:397: There is a wand
/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/pytorch_lightning/callbacks/model_checkpoint.py:654: Ch

```

	Name	Type	Params	Mode
0	model	BestModel	61.3 K	train
1	criterion	BCELoss	0	eval

```

61.3 K Trainable params
0 Non-trainable params
61.3 K Total params
0.245 Total estimated model params size (MB)
6 Modules in train mode
1 Modules in eval mode

```

```

start training
Sanity Checking DataLoader 0: 0%|          | 0/2 [00:00<?, ?it/s]/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/si
/Users/jirayuwat/anaconda3/envs/nlp/lib/pytho
Epoch 2: 100%|██████████| 32152/32152 [09:55<00:00, 54.02it/s, v_num=04og]\Trainer.fit` stopped: `max_epochs=3` reached.
Epoch 2: 100%|██████████| 32152/32152 [09:55<00:00, 54.02it/s, v_num=04og]

```



```

1 evaluate(test_loader, my_best_model.to('mps:0'))

```

```

100%|██████████| 4438/4438 [00:38<00:00, 115.39it/s]
{'accuracy': 0.9864322194341497,
 'f1_score': 0.9732915180815483,
 'precision': 0.9633137650231901,
 'recall': 0.9834781283624118}

```