

✓ HOMEWORK 6: TEXT CLASSIFICATION

In this homework, you will create models to classify texts from TRUE call-center. There are two classification tasks:

1. Action Classification: Identify which action the customer would like to take (e.g. enquire, report, cancel)
2. Object Classification: Identify which object the customer is referring to (e.g. payment, true money, internet, roaming)

We will focus only on the Object Classification task for this homework.

In this homework, you are asked to compare different text classification models in terms of accuracy and inference time.

You will need to build 3 different models.

1. A model based on tf-idf
2. A model based on MUSE
3. A model based on wangchanBERTa

You will be asked to submit 3 different files (.pdf from .ipynb) that do the 3 different models. Finally, answer the accuracy and runtime numbers in MCV.

This homework is quite free form, and your answer may vary. We hope that the processing during the course of this assignment will make you think more about the design choices in text classification.

```
1 !wget --no-check-certificate https://www.dropbox.com/s/37u83g55p19kvrl/clean-phone-data-for-students.csv -O ./clean-phone-
--2025-02-15 15:16:27-- https://www.dropbox.com/s/37u83g55p19kvrl/clean-phone-data-for-students.csv
Resolving www.dropbox.com (www.dropbox.com)... 2620:100:6035:18::a27d:5512, 162.125.85.18
Connecting to www.dropbox.com (www.dropbox.com)|2620:100:6035:18::a27d:5512|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://www.dropbox.com/scl/fi/8h8hvs9uj6o0524lfe4i/clean-phone-data-for-students.csv?rlkey=lwv5xbf16jerehny3
--2025-02-15 15:16:27-- https://www.dropbox.com/scl/fi/8h8hvs9uj6o0524lfe4i/clean-phone-data-for-students.csv?rlkey=lw
Reusing existing connection to [www.dropbox.com]:443.
HTTP request sent, awaiting response... 302 Found
Location: https://ucce80429c1d534064b3a6d9f042.dl.dropboxusercontent.com/cd/0/inline/CkLrInUX8rJuXukZ2pm0zA1HWeZom4RNG8z
--2025-02-15 15:16:28-- https://ucce80429c1d534064b3a6d9f042.dl.dropboxusercontent.com/cd/0/inline/CkLrInUX8rJuXukZ2pm0
Resolving ucce80429c1d534064b3a6d9f042.dl.dropboxusercontent.com (ucce80429c1d534064b3a6d9f042.dl.dropboxusercontent.com)
Connecting to ucce80429c1d534064b3a6d9f042.dl.dropboxusercontent.com (ucce80429c1d534064b3a6d9f042.dl.dropboxusercontent
HTTP request sent, awaiting response... 200 OK
Length: 2518977 (2.4M) [text/plain]
Saving to: './clean-phone-data-for-students.csv'

./clean-phone-data- 100%[=====>] 2.40M 4.07MB/s in 0.6s

2025-02-15 15:16:30 (4.07 MB/s) - './clean-phone-data-for-students.csv' saved [2518977/2518977]

1 !pip install -q pythainlp
```

✓ Import Libs

```
1 %matplotlib inline
2 import pandas
3 import sklearn
4 import numpy as np
5 import time
6 import matplotlib.pyplot as plt
7 import pandas as pd
8 from pprint import pprint
9
10 from torch.utils.data import Dataset
11 from IPython.display import display
12 from collections import defaultdict
13 from sklearn.metrics import accuracy_score
```

✓ Loading data

First, we load the data from disk into a Dataframe.

A Dataframe is essentially a table, or 2D-array/Matrix with a name for each column.


```
1 data_df = pd.read_csv('clean-phone-data-for-students.csv')
```

Let's preview the data.

```

1 # Show the top 5 rows
2 display(data_df.head())
3 # Summarize the data
4 data_df.describe()

```



	Sentence Utterance	Action	Object
0	<PHONE_NUMBER_REMOVED> ผมไปจ่ายเงินที่ Counte...	enquire	payment
1	internet ยังความเร็วอยู่เท่าไรครับ	enquire	package
2	ตะกี้ไปชำระค่าบริการไปแล้ว แต่ยังไม่ใช้งานไม่ได้...	report	suspend
3	พี่คะยัง ใช้ internet ไม่ได้เลยคะ เป็นเครือ...	enquire	internet
4	ฮาโหล ค่ะ พอตัวเมื่อวานเปิดซิมทรูมูฟ แต่มั่น...	report	phone_issues

	Sentence Utterance	Action	Object
count	16175	16175	16175
unique	13389	10	33
top	บริการอื่นๆ	enquire	service
freq	97	10377	2525

✓ Data cleaning

We call the `DataFrame.describe()` again. Notice that there are 33 unique labels/classes for object and 10 unique labels for action that the model will try to predict. But there are unwanted duplications e.g. `Idd,idd,lotalty_card,Lotalty_card`

Also note that, there are 13389 unqiue sentence utterances from 16175 utterances. You have to clean that too!


#TODO 0.1:

You will have to remove unwanted label duplications as well as duplications in text inputs. Also, you will have to trim out unwanted whitespaces from the text inputs. This shouldn't be too hard, as you have already seen it in the demo.

```

1 display(data_df.describe())
2 display(data_df.Object.unique())
3 display(data_df.Action.unique())

```



	Sentence Utterance	Action	Object
count	16175	16175	16175
unique	13389	10	33
top	บริการอื่นๆ	enquire	service
freq	97	10377	2525


```

array(['payment', 'package', 'suspend', 'internet', 'phone_issues',
       'service', 'nonTrueMove', 'balance', 'detail', 'bill', 'credit',
       'promotion', 'mobile_setting', 'iservice', 'roaming', 'truemoney',
       'information', 'lost_stolen', 'balance_minutes', 'idd',
       'TrueMoney', 'garbage', 'Payment', 'IDD', 'ringtone', 'Idd',
       'rate', 'loyalty_card', 'contact', 'officer', 'Balance', 'Service',
       'Loyalty_card'], dtype=object)
array(['enquire', 'report', 'cancel', 'Enquire', 'buy', 'activate',
       'request', 'Report', 'garbage', 'change'], dtype=object)

```

```

1 clean_data_time = time.time()
2
3 # Group the duplicate label
4 data_df.dropna(subset=['Object'], inplace=True)
5 data_df['Object'] = data_df['Object'].apply(lambda x: x.lower())
6
7 # Clean the data
8 data_df['Sentence Utterance'] = data_df['Sentence Utterance'].apply(lambda x: str(x).strip())
9 data_df['Sentence Utterance'] = data_df['Sentence Utterance'].apply(lambda x: x.lower())
10 data_df.drop_duplicates(subset=['Sentence Utterance'], inplace=True)
11
12 # Drop the unused columns
13 data_df.drop(columns=['Action'], inplace=True)
14
15 clean_data_time = time.time() - clean_data_time

```

Split data into train, validation, and test sets (normally the ratio will be 80:10:10 , respectively). We recommend to use `train_test_spilt` from `scikit-learn` to split the data into train, validation, test set.

In addition, it should split the data that distribution of the labels in train, validation, test set are similar. There is **stratify** option to handle this issue.

https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html

Make sure the same data splitting is used for all models.

```
1 from sklearn.model_selection import train_test_split
2 from collections import Counter
3
4 split_data_time = time.time()
5
6 # For the object column, we will only keep the object that has more than 2% of the total data
7 object_counter = Counter(data_df['Object'])
8 stratify_col = data_df['Object'].apply(lambda x: 'other' if object_counter[x] < 0.02*len(data_df) else x)
9 train_df, test_df = train_test_split(data_df, test_size=0.2, random_state=4242, stratify=stratify_col)
10
11 object_counter = Counter(test_df['Object'])
12 stratify_col = test_df['Object'].apply(lambda x: 'other' if object_counter[x] < 0.02*len(test_df) else x)
13 test_df, val_df = train_test_split(test_df, test_size=0.5, random_state=4242, stratify=stratify_col)
14
15 train_df = train_df.reset_index(drop=True)
16 test_df = test_df.reset_index(drop=True)
17 val_df = val_df.reset_index(drop=True)
18
19 print(f"Train size: {len(train_df)}")
20 print(f"Test size: {len(test_df)}")
21 print(f"Val size: {len(val_df)}")
22
23 split_data_time = time.time() - split_data_time
```

```
➡ Train size: 10689
   Test size: 1336
   Val size: 1337
```

```
1 # Save the data
2 # train_df.to_csv('train.csv', index=False)
3 # test_df.to_csv('test.csv', index=False)
4 # val_df.to_csv('val.csv', index=False)
```

✓ Model 2 MUSE

Build a simple logistic regression model using features from the MUSE model.

Which MUSE model will you use? Why?

Ans: sentence-transformers/use-cmlm-multilingual because it is easy to use and was trained on many languages(100+). The community also has a lot of good feedback on this model. (about 3,000 downloads last month)

MUSE is typically used with tensorflow. However, there are some pytorch conversions made by some people.

<https://huggingface.co/sentence-transformers/use-cmlm-multilingual> <https://huggingface.co/dayyass/universal-sentence-encoder-multilingual-large-3-pytorch>

```
1 !pip install -qU sentence-transformers
```

```
1 import pandas as pd
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.metrics import classification_report, accuracy_score
4 from sentence_transformers import SentenceTransformer
```

```
1 train_df, test_df, val_df = pd.read_csv('train.csv'), pd.read_csv('test.csv'), pd.read_csv('val.csv')
```

```
1 encoder_model = SentenceTransformer('sentence-transformers/use-cmlm-multilingual')
```

➡ Some weights of the model checkpoint at sentence-transformers/use-cmlm-multilingual were not used when initializing Bert
- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with anot
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly id

```
1 encoder_time = time.time()
2
3 train_x = encoder_model.encode(train_df['Sentence Utterance'].values, show_progress_bar=True)
4 test_x = encoder_model.encode(test_df['Sentence Utterance'].values, show_progress_bar=True)
5 val_x = encoder_model.encode(val_df['Sentence Utterance'].values, show_progress_bar=True)
6
7 train_y = train_df['Object']
```

```

8 test_y = test_df['Object']
9 val_y = val_df['Object']
10
11 encoder_time = time.time() - encoder_time

→ Batches: 0%|          | 0/335 [00:00<?, ?it/s]
   Batches: 0%|          | 0/42 [00:00<?, ?it/s]
   Batches: 0%|          | 0/42 [00:00<?, ?it/s]

1 train_time = time.time()
2
3 model = LogisticRegression(max_iter=1000, penalty='elasticnet', solver='saga', l1_ratio=0.5, C=1.0, n_jobs=-1)
4 model.fit(train_x, train_y)
5
6 train_time = time.time() - train_time

```

```

1 all_inference_time = time.time()
2
3 # Predict the data
4 train_pred = model.predict(train_x)
5 test_pred = model.predict(test_x)
6 val_pred = model.predict(val_x)
7
8 all_inference_time = time.time() - all_inference_time
9
10 # Calculate the accuracy
11 train_acc = accuracy_score(train_y, train_pred)
12 test_acc = accuracy_score(test_y, test_pred)
13 val_acc = accuracy_score(val_y, val_pred)
14
15 print(f"Train accuracy: {train_acc}")
16 print(f"Test accuracy: {test_acc}")
17 print(f"Val accuracy: {val_acc}")

```

```

→ Train accuracy: 0.7254186546917392
   Test accuracy: 0.7148203592814372
   Val accuracy: 0.7120418848167539

```

```

1 # Print the classification report
2 print("Val classification report")
3 print(classification_report(val_y, val_pred))
4
5 import pickle
6
7 # Save the classificaion report
8 classification_report_dict = classification_report(val_y, val_pred, output_dict=True)
9 with open('classification_report_muse.pkl', 'wb') as f:
10     pickle.dump(classification_report_dict, f)

```

```

→ Val classification report
      precision    recall  f1-score   support

 balance      0.73      0.83      0.78      148
balance_minutes      0.00      0.00      0.00         5
      bill      0.62      0.52      0.57         54
      credit      1.00      0.65      0.79         17
      detail      0.50      0.27      0.35         33
      garbage      0.00      0.00      0.00          6
      idd      0.91      0.71      0.80         14
      information      0.64      0.48      0.55         29
      internet      0.69      0.83      0.76      179
      iservice      0.00      0.00      0.00          3
      lost_stolen      0.96      0.77      0.85         30
      loyalty_card      1.00      0.25      0.40          8
      mobile_setting      0.57      0.14      0.23         28
      nontruemove      0.60      0.14      0.23         21
      officer      0.00      0.00      0.00          1
      package      0.65      0.75      0.70      179
      payment      0.69      0.72      0.70         64
      phone_issues      0.61      0.67      0.64         58
      promotion      0.72      0.70      0.71      115
      rate      0.00      0.00      0.00          6
      ringtone      0.88      0.64      0.74         11
      roaming      0.76      0.76      0.76         17
      service      0.77      0.84      0.80      211
      suspend      0.73      0.77      0.75         73
      truemoney      0.92      0.85      0.88         27

 accuracy      0.60      0.49      0.52      1337
  macro avg      0.60      0.49      0.52      1337
  weighted avg      0.70      0.71      0.70      1337

```

```
/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined: ROC curve estimated from 0 true positives
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined: ROC curve estimated from 0 true positives
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined: ROC curve estimated from 0 true positives
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined: ROC curve estimated from 0 true positives
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/Users/jirayuwat/anaconda3/envs/nlp/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined: ROC curve estimated from 0 true positives
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
1 print(f'''
2 All preprocessing time: {clean_data_time + split_data_time + encoder_time:.2f} seconds
3 - Clean data time: {clean_data_time:.2f} seconds
4 - Split data time: {split_data_time:.2f} seconds
5 - Encoder time: {encoder_time:.2f} seconds
6 Training time: {train_time:.2f} seconds
7 Inference time: {all_inference_time:.2f} seconds
8 '''.strip())
```

```
➡ All preprocessing time: 45.99 seconds
  - Clean data time: 0.01 seconds
  - Split data time: 0.01 seconds
  - Encoder time: 45.96 seconds
Training time: 24.70 seconds
Inference time: 0.01 seconds
```