

✓ HOMEWORK 6: TEXT CLASSIFICATION

In this homework, you will create models to classify texts from TRUE call-center. There are two classification tasks:

1. Action Classification: Identify which action the customer would like to take (e.g. enquire, report, cancel)
2. Object Classification: Identify which object the customer is referring to (e.g. payment, true money, internet, roaming)

We will focus only on the Object Classification task for this homework.

In this homework, you are asked to compare different text classification models in terms of accuracy and inference time.

You will need to build 3 different models.

1. A model based on tf-idf
2. A model based on MUSE
3. A model based on wangchanBERTa

You will be asked to submit 3 different files (.pdf from .ipynb) that do the 3 different models. Finally, answer the accuracy and runtime numbers in MCV.

This homework is quite free form, and your answer may vary. We hope that the processing during the course of this assignment will make you think more about the design choices in text classification.

```
1 !wget --no-check-certificate https://www.dropbox.com/s/37u83g55p19kvrl/clean-phone-data-for-students.csv -O ./clean-phone
--2025-02-15 09:30:32-- https://www.dropbox.com/s/37u83g55p19kvrl/clean-phone-data-for-students.csv
Resolving www.dropbox.com (www.dropbox.com)... 162.125.2.18, 2620:100:6017:18::a27d:212
Connecting to www.dropbox.com (www.dropbox.com)|162.125.2.18|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://www.dropbox.com/scl/fi/8h8hvs9uj6o0524lfe4i/clean-phone-data-for-students.csv?rlkey=lwv5xbf16jerehny3
--2025-02-15 09:30:32-- https://www.dropbox.com/scl/fi/8h8hvs9uj6o0524lfe4i/clean-phone-data-for-students.csv?rlkey=lw
Reusing existing connection to www.dropbox.com:443.
HTTP request sent, awaiting response... 302 Found
Location: https://ucdf0b5923cd38339989bc62a9f.dl.dropboxusercontent.com/cd/0/inline/CkJvDe1QdFc6hKiyZjSdzab3ahggm-dltXL
--2025-02-15 09:30:32-- https://ucdf0b5923cd38339989bc62a9f.dl.dropboxusercontent.com/cd/0/inline/CkJvDe1QdFc6hKiyZjSd
Resolving ucdf0b5923cd38339989bc62a9f.dl.dropboxusercontent.com (ucdf0b5923cd38339989bc62a9f.dl.dropboxusercontent.com)
Connecting to ucdf0b5923cd38339989bc62a9f.dl.dropboxusercontent.com (ucdf0b5923cd38339989bc62a9f.dl.dropboxusercontent
HTTP request sent, awaiting response... 200 OK
Length: 2518977 (2.4M) [text/plain]
Saving to: './clean-phone-data-for-students.csv'

./clean-phone-data- 100%[=====>] 2.40M --.-KB/s in 0.05s

2025-02-15 09:30:33 (46.7 MB/s) - './clean-phone-data-for-students.csv' saved [2518977/2518977]
```

```
1 !pip install -q pythainlp
```

✓ Import Libs

```
1 %matplotlib inline
2 import pandas
3 import sklearn
4 import numpy as np
5 import time
6 import matplotlib.pyplot as plt
7 import pandas as pd
8 from pprint import pprint
9
10 from torch.utils.data import Dataset
11 from IPython.display import display
12 from collections import defaultdict
13 from sklearn.metrics import accuracy_score
```

✓ Loading data

First, we load the data from disk into a Dataframe.

A Dataframe is essentially a table, or 2D-array/Matrix with a name for each column.

```
1 data_df = pd.read_csv('clean-phone-data-for-students.csv')
```

Let's preview the data.

```

1 # Show the top 5 rows
2 display(data_df.head())
3 # Summarize the data
4 data_df.describe()

```

	Sentence Utterance	Action	Object
0	<PHONE_NUMBER_REMOVED> ผมไปจ่ายเงินที่ Counte...	enquire	payment
1	internet ยังความเร็วอยู่เท่าไร ครับ	enquire	package
2	ตะกี้ไปชำระค่าบริการไปแล้ว แต่ยังไม่ใช้งานไม่ได้...	report	suspend
3	พี่คะยัง ใช้ internet ไม่ได้เลยคะ เป็นเครือ...	enquire	internet
4	ฮาโหล ค่ะ พอตัวเมื่อวานเปิดซิมทรูมูฟ แต่มั่น...	report	phone_issues

	Sentence Utterance	Action	Object
count	16175	16175	16175
unique	13389	10	33
top	บริการอื่นๆ	enquire	service
freq	97	10377	2525

✓ Data cleaning

We call the `DataFrame.describe()` again. Notice that there are 33 unique labels/classes for object and 10 unique labels for action that the model will try to predict. But there are unwanted duplications e.g. `Idd,idd,lotalty_card,Lotalty_card`

Also note that, there are 13389 unique sentence utterances from 16175 utterances. You have to clean that too!

#TODO 0.1:

You will have to remove unwanted label duplications as well as duplications in text inputs. Also, you will have to trim out unwanted whitespaces from the text inputs. This shouldn't be too hard, as you have already seen it in the demo.

```

1 display(data_df.describe())
2 display(data_df.Object.unique())
3 display(data_df.Action.unique())

```

	Sentence Utterance	Action	Object
count	16175	16175	16175
unique	13389	10	33
top	บริการอื่นๆ	enquire	service
freq	97	10377	2525


```

array(['payment', 'package', 'suspend', 'internet', 'phone_issues',
       'service', 'nonTrueMove', 'balance', 'detail', 'bill', 'credit',
       'promotion', 'mobile_setting', 'iservice', 'roaming', 'truemoney',
       'information', 'lost_stolen', 'balance_minutes', 'idd',
       'TrueMoney', 'garbage', 'Payment', 'IDD', 'ringtone', 'Idd',
       'rate', 'loyalty_card', 'contact', 'officer', 'Balance', 'Service',
       'Loyalty_card'], dtype=object)
array(['enquire', 'report', 'cancel', 'Enquire', 'buy', 'activate',
       'request', 'Report', 'garbage', 'change'], dtype=object)

```

```

1 clean_data_time = time.time()
2
3 # Group the duplicate label
4 data_df.dropna(subset=['Object'], inplace=True)
5 data_df['Object'] = data_df['Object'].apply(lambda x: x.lower())
6
7 # Clean the data
8 data_df['Sentence Utterance'] = data_df['Sentence Utterance'].apply(lambda x: str(x).strip())
9 data_df['Sentence Utterance'] = data_df['Sentence Utterance'].apply(lambda x: x.lower())
10 data_df.drop_duplicates(subset=['Sentence Utterance'], inplace=True)
11
12 # Drop the unused columns
13 data_df.drop(columns=['Action'], inplace=True)
14
15 clean_data_time = time.time() - clean_data_time

```

Split data into train, validation, and test sets (normally the ratio will be 80:10:10, respectively). We recommend to use `train_test_split` from `scikit-learn` to split the data into train, validation, test set.

In addition, it should split the data that distribution of the labels in train, validation, test set are similar. There is **stratify** option to handle this issue.

https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html

Make sure the same data splitting is used for all models.

```
1 from sklearn.model_selection import train_test_split
2 from collections import Counter
3
4 split_data_time = time.time()
5
6 # For the object column, we will only keep the object that has more than 2% of the total data
7 object_counter = Counter(data_df['Object'])
8 stratify_col = data_df['Object'].apply(lambda x: 'other' if object_counter[x] < 0.02*len(data_df) else x)
9 train_df, test_df = train_test_split(data_df, test_size=0.2, random_state=4242, stratify=stratify_col)
10
11 object_counter = Counter(test_df['Object'])
12 stratify_col = test_df['Object'].apply(lambda x: 'other' if object_counter[x] < 0.02*len(test_df) else x)
13 test_df, val_df = train_test_split(test_df, test_size=0.5, random_state=4242, stratify=stratify_col)
14
15 train_df = train_df.reset_index(drop=True)
16 test_df = test_df.reset_index(drop=True)
17 val_df = val_df.reset_index(drop=True)
18
19 print(f"Train size: {len(train_df)}")
20 print(f"Test size: {len(test_df)}")
21 print(f"Val size: {len(val_df)}")
22
23 split_data_time = time.time() - split_data_time
```

```
➡ Train size: 10689
   Test size: 1336
   Val size: 1337
```

```
1 # Save the data
2 train_df.to_csv('train.csv', index=False)
3 test_df.to_csv('test.csv', index=False)
4 val_df.to_csv('val.csv', index=False)
```

✓ Model 3 WangchanBERTa

We ask you to train a WangchanBERTa-based model.

We recommend you use the thaixtransformers fork (which we used in the PoS homework). <https://github.com/PyThaiNLP/thaixtransformers>

The structure of the code will be very similar to the PoS homework. You will also find the huggingface [tutorial](#) useful. Or you can also add a softmax layer by yourself just like in the previous homework.

Which WangchanBERTa model will you use? Why? (Don't forget to clean your text accordingly).

Ans:

```
1 %pip install -q wandb
2 %pip install -q transformers==4.30.1 datasets evaluate thaixtransformers
3 %pip install -q emoji pythainlp sefr_cut tinydb sequeval sentencepiece pydantic jsonlines
4 %pip install -q peft==0.10.0

1 import pandas as pd
2 from thaixtransformers import Tokenizer
3 from transformers import AutoModelForSequenceClassification, DataCollatorWithPadding, TrainingArguments, Trainer
4 from datasets import Dataset
5 import time
6 from sklearn.metrics import accuracy_score

1 create_dataset_time = time.time()
2
3 train_df, test_df, val_df = pd.read_csv('train.csv'), pd.read_csv('test.csv'), pd.read_csv('val.csv')
4 train_df.columns = ['text', 'label']
5 test_df.columns = ['text', 'label']
6 val_df.columns = ['text', 'label']
7
8 label2id = {label: i for i, label in enumerate(sorted(train_df['label'].unique()))}
9 id2label = {i: label for label, i in label2id.items()}
10 train_df['label'] = train_df['label'].apply(lambda x: label2id[x])
11 test_df['label'] = test_df['label'].apply(lambda x: label2id[x])
12 val_df['label'] = val_df['label'].apply(lambda x: label2id[x])
13
```

```

14 # Create dataset
15 train_dataset = Dataset.from_pandas(train_df)
16 test_dataset = Dataset.from_pandas(test_df)
17 val_dataset = Dataset.from_pandas(val_df)

1 tokenizer = Tokenizer("airesearch/wangchanberta-base-wiki-newmm")
2 model = AutoModelForSequenceClassification.from_pretrained("airesearch/wangchanberta-base-wiki-newmm",
3                                                         num_labels=train_df['label'].max()+1,
4                                                         label2id=label2id,
5                                                         id2label=id2label)
6
7 data_collator = DataCollatorWithPadding(tokenizer=tokenizer)

🔄 /usr/local/lib/python3.11/dist-packages/huggingface_hub/file_download.py:795: FutureWarning: `resume_download` is deprecated and will be removed in version 0.11; the automatic resumption of downloads will be the default, please read more at https://huggingface.co/docs/huggingface_hub/CHANGELOG.html#0.11
warnings.warn(
The tokenizer class you load from this checkpoint is not the same type as the class this function is called from. It may be the same class but a different version.
The tokenizer class you load from this checkpoint is 'RobertaTokenizer'.
The class this function is called from is 'ThaiWordsNewmmTokenizer'.
/usr/local/lib/python3.11/dist-packages/transformers/modeling_utils.py:463: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value), which is the unsafe legacy default behavior. It is recommended to set `weights_only=True` to load the checkpoint safely. `weights_only=True` tells the module that we expect only state-dict tensors and not tensors with backends.
return torch.load(checkpoint_file, map_location="cpu")
Some weights of the model checkpoint at airesearch/wangchanberta-base-wiki-newmm were not used when initializing RobertaForSequenceClassification:
- This IS expected if you are initializing RobertaForSequenceClassification from the checkpoint of a model trained on an
- This IS NOT expected if you are initializing RobertaForSequenceClassification from the checkpoint of a model that you
Some weights of RobertaForSequenceClassification were not initialized from the model checkpoint at airesearch/wangchanberta-base-wiki-newmm and are newly initialized from the current random state. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

1 # Tokenize the data
2 train_dataset = train_dataset.map(lambda x: tokenizer(x['text'], padding="max_length", truncation=True), batched=True)
3 test_dataset = test_dataset.map(lambda x: tokenizer(x['text'], padding="max_length", truncation=True), batched=True)
4 val_dataset = val_dataset.map(lambda x: tokenizer(x['text'], padding="max_length", truncation=True), batched=True)
5
6 create_dataset_time = time.time() - create_dataset_time

🔄 Map: 100% 10689/10689 [00:09<00:00, 1229.44 examples/s]
Asking to pad to max_length but no maximum length is provided and the model has no predefined maximum length. Default to 512.
Asking to truncate to max_length but no maximum length is provided and the model has no predefined maximum length. Default to 512.
Map: 100% 1336/1336 [00:01<00:00, 875.73 examples/s]
Map: 100% 1337/1337 [00:01<00:00, 1036.96 examples/s]

1 def compute_metrics(pred):
2     labels = pred.label_ids
3     preds = pred.predictions.argmax(-1)
4     return {"accuracy": accuracy_score(labels, preds)}
5
6 training_args = TrainingArguments(
7     #####
8     output_dir="text_classification",
9     learning_rate=2e-5,
10    num_train_epochs=2,
11    weight_decay=0.01,
12    push_to_hub=False
13    #####
14 )
15
16 trainer = Trainer(
17     #####
18     model=model,
19     args=training_args,
20     train_dataset=train_dataset,
21     eval_dataset=test_dataset,
22     data_collator=data_collator,
23     compute_metrics=compute_metrics,
24     #####
25 )
26
27 train_time = time.time()
28 trainer.train()
29 train_time = time.time() - train_time


```

 /usr/local/lib/python3.11/dist-packages/transformers/optimization.py:411: FutureWarning: This implementation of AdamW is warnings.warn(
[2674/2674 08:25, Epoch 2/2]


Step Training Loss

500	1.898900
1000	1.184500
1500	1.021700
2000	0.759500
2500	0.766700

```
1 all_inference_time = time.time()
2
3 # Predict the data
4 train_pred = trainer.predict(train_dataset)
5 test_pred = trainer.predict(test_dataset)
6 val_pred = trainer.predict(val_dataset)
7
8 all_inference_time = time.time() - all_inference_time
9
10 # Calculate the accuracy
11 train_acc = train_pred.metrics['test_accuracy']
12 test_acc = test_pred.metrics['test_accuracy']
13 val_acc = val_pred.metrics['test_accuracy']
14
15 print(f"Train accuracy: {train_acc}")
16 print(f"Test accuracy: {test_acc}")
17 print(f"Val accuracy: {val_acc}")
```

 Train accuracy: 0.8529329216952006
Test accuracy: 0.7574850299401198
Val accuracy: 0.7471952131637996

```
1 from sklearn.metrics import classification_report, accuracy_score
2
3 # Print the classification report
4 print("Val classification report")
5 print(classification_report(val_pred.label_ids, val_pred.predictions.argmax(-1)))
6
7 import pickle
8
9 # Save the classificaion report
10 classification_report_dict = classification_report(val_pred.label_ids, val_pred.predictions.argmax(-1), output_dict=True)
11 with open('classification_report_wangchanberta.pkl', 'wb') as f:
12     pickle.dump(classification_report_dict, f)
```

 Val classification report

	precision	recall	f1-score	support
0	0.82	0.80	0.81	148
1	0.00	0.00	0.00	5
2	0.61	0.63	0.62	54
4	1.00	0.88	0.94	17
5	0.57	0.24	0.34	33
6	0.00	0.00	0.00	6
7	0.75	0.64	0.69	14
8	0.62	0.55	0.58	29
9	0.76	0.79	0.77	179
10	0.00	0.00	0.00	3
11	1.00	0.93	0.97	30
12	1.00	0.38	0.55	8
13	0.46	0.43	0.44	28
14	1.00	0.14	0.25	21
15	0.00	0.00	0.00	1
16	0.71	0.82	0.76	179
17	0.65	0.75	0.70	64
18	0.69	0.72	0.71	58
19	0.74	0.70	0.72	115
20	0.00	0.00	0.00	6
21	0.92	1.00	0.96	11
22	0.67	0.71	0.69	17
23	0.79	0.86	0.82	211
24	0.78	0.86	0.82	73
25	0.90	0.96	0.93	27
accuracy			0.75	1337
macro avg	0.62	0.55	0.56	1337
weighted avg	0.74	0.75	0.73	1337

```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is il
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is il
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is il
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is il
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is il
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is il
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

```

1 print(f'''
2 All preprocessing time: {clean_data_time + split_data_time + create_dataset_time:.2f} seconds
3 - Clean data time: {clean_data_time:.2f} seconds
4 - Split data time: {split_data_time:.2f} seconds
5 - Create dataset time: {create_dataset_time:.2f} seconds
6 Training time: {train_time:.2f} seconds
7 Inference time: {all_inference_time:.2f} seconds
8 '''.strip())

```

```

🔗 All preprocessing time: 20.37 seconds
  - Clean data time: 0.03 seconds
  - Split data time: 0.05 seconds
  - Create dataset time: 20.29 seconds
Training time: 507.56 seconds
Inference time: 32.75 seconds

```

After you

Comparison

After you have completed the 3 models, compare the accuracy, ease of implementation, and inference speed (from cleaning, tokenization, till model compute) between the three models in mycourseville.